

BASiS: Adapting Conversational Behavior in Purpose-Specific Dialogue Systems through Bayesian Data Selection

ANONYMOUS AUTHOR(S)

We propose BASiS (Bayesian Style-based Selection), a method for adapting a local large language model (LLM) to the conversational styles required in purpose-specific dialogue through Bayesian data selection from small, high-quality corpora. Although recent LLMs can generate fluent and accurate responses, settings such as emotional support, mathematics tutoring, and medical interviewing require not only correct content but also an appropriate way of responding. Such styles are often captured in small dialogue corpora that are too limited for direct LLM training. Large general-purpose dialogue corpora provide sufficient scale, yet may contain too few examples that match the target style. BASiS extracts dialogue states, response strategies, and state transitions from a small corpus and represents them as a Bayesian dialogue model. It then scores dialogues in a large corpus according to their compatibility with the target style, selects training data, constructs preference pairs, and trains the LLM using Direct Preference Optimization (DPO). Its key advantage is that it converts the implicit conversational dynamics of a small corpus into a criterion for identifying suitable data in a much larger corpus. This reduces the burden of manually designing detailed dialogue guidelines while allowing selection to reflect response timing and multi-turn progression, which generic quality or topical similarity alone cannot capture. By maintaining dialogue states as a probability distribution rather than committing to a single response strategy, BASiS can preserve multiple valid strategies for the same context. We evaluated BASiS in emotional support, mathematics tutoring, and medical interviewing by comparing the base model, a model trained on BASiS-selected data, and a model trained on randomly selected in-category data. Across all three settings, the BASiS-trained model achieved the highest scores for reproducing the characteristic response policies and conversational progression of the target corpora.

CCS Concepts: • **Computing methodologies** → **Natural language generation**; **Discourse, dialogue and pragmatics**; Learning from implicit feedback; • **Human-centered computing** → *Natural language interfaces*.

Additional Key Words and Phrases: purpose-specific dialogue systems, conversational style, data selection, Bayesian dialogue modeling, preference optimization, emotional support dialogue, large language models

ACM Reference Format:

Anonymous Author(s). 2026. BASiS: Adapting Conversational Behavior in Purpose-Specific Dialogue Systems through Bayesian Data Selection. In *Proceedings of 32nd Annual ACM Conference on Intelligent User Interfaces (IUI '27)*. ACM, New York, NY, USA, 17 pages.

1 Introduction

Advances in instruction tuning and reinforcement learning from human feedback (RLHF) have enabled recent large language models (LLMs) to follow instructions, generate accurate responses, and sustain natural and coherent multi-turn conversations across a wide range of topics [11]. In purpose-specific dialogue such as emotional support, mathematics tutoring, and medical history-taking, however, a system must do more than provide correct information: it must choose an appropriate response policy based on the interlocutor's state and the current stage of the dialogue [7, 13]. For example, it may need to acknowledge a person's feelings before offering advice, encourage a learner to reason independently,

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

or elicit necessary symptoms in an appropriate order. Such conversational styles are reflected in small, high-quality corpora developed for specific purposes, but these corpora are limited in size for direct use in LLM training. Large general-purpose corpora offer greater volume and diversity, yet they do not necessarily include enough dialogues that match the target style. A method is therefore needed to use the style captured in a small, high-quality corpus to select suitable dialogue data from a large general-purpose corpus.

Several approaches have been explored to generate responses suited to purpose-specific dialogue. Guideline-based response control explicitly defines a desirable dialogue policy in advance and guides response generation to follow it. GuidedTOD, proposed by Wen et al., probabilistically represents the flow of task-oriented dialogue and reranks candidate responses while updating a guideline-based policy as the dialogue progresses; it achieved higher guideline compliance than previous methods [16]. Another line of work predicts a response strategy from the current user state. EmoDynamix, proposed by Wan et al., models the relationships between multiple user emotion states and response strategies as a heterogeneous graph, achieving strong strategy-prediction performance while making the selection process interpretable [15]. Research on selecting LLM training data has further shown that carefully chosen, high-quality data can efficiently adapt model behavior even when the selected dataset is small [18, 21].

These approaches nevertheless leave three challenges. First, guideline-based response control requires desirable dialogue policies to be written manually for each purpose or domain. Conversational style is often expressed implicitly through tone, response timing, and progression across multiple turns, making it difficult to describe comprehensively as explicit rules. Second, response-strategy prediction generally predicts a single next strategy for the current state. In real dialogue, however, several strategies may be valid in the same context, such as continuing to empathize, asking a question to explore the issue, or organizing the available information. A single prediction cannot adequately represent the diversity of responses found in a target corpus. Third, many existing data-selection methods focus on general qualities such as accuracy, usefulness, diversity, and task coverage. Even when two responses are comparable on these dimensions, they may differ in when they are given and in how they change the conversational state. Unless dialogue history and transitions between response strategies are included in the selection criterion, it is difficult to reliably select data that preserve the characteristic dialogue flow of the target corpus.

To address these challenges, we propose BASiS (Bayesian Style-based Selection), a method that adapts a local LLM to the conversational style required for purpose-specific dialogue through Bayesian data selection based on a small, high-quality corpus. BASiS first uses an LLM to analyze the small corpus and structures each dialogue in terms of conversational states, which describe the interlocutor’s situation or the stage of the dialogue; response strategies, which describe response policies such as empathy, questioning, and advice; and state transitions caused by the responses. It then constructs a Bayesian dialogue model from this structure, using transition probabilities between states and observation likelihoods for response strategies in each state. The model scores each dialogue instance in a large general-purpose corpus according to its compatibility with the target conversational style and selects data for training. Finally, BASiS constructs preference pairs from the selected data and applies Direct Preference Optimization (DPO) with Low-Rank Adaptation (LoRA) to a local LLM, thereby adapting the model to the style captured in the small corpus.

A central feature of BASiS is that it represents conversational style not as individual response expressions or detailed, manually authored guidelines, but as probabilistic transitions between conversational states and response strategies. This allows the implicit dialogue flow in the corpus to guide data selection without requiring detailed guidelines for each purpose. Rather than treating a single response strategy as the only correct choice, BASiS maintains a probability distribution over conversational states and evaluates candidate responses through Bayesian updating. It can therefore accommodate conversational diversity in which multiple strategies may be valid for the same context. Moreover, the

same procedure can be applied to different purpose-specific settings, including emotional support, mathematics tutoring, and medical history-taking, by changing the small reference corpus and the large corpus from which data are selected. Using the selected data for LoRA-based DPO allows the target style to be incorporated while retaining the general dialogue capabilities of the local LLM.

The remainder of this paper is organized as follows. Section 2 reviews related work on guideline-based response control, dialogue-state-based response-strategy prediction, and training-data selection for LLMs. Section 3 describes the overall architecture of BASiS and the role of each module. Section 4 presents the experiments used to evaluate BASiS, their results, and our discussion. Finally, Section 5 concludes the paper.

2 Related Work

One approach to realizing purpose-specific behavior in dialogue systems is to provide desirable response policies as natural-language guidelines and control response generation accordingly. Gupta et al. proposed DialGuide, which describes in natural language both the conversational contexts in which a guideline applies and the content that a response should include [4]. DialGuide addresses guideline selection, guideline-grounded response generation, and verification that generated responses satisfy the selected guideline. Experiments on open-domain and safety-related dialogue data showed that it can generate safe and engaging responses that follow developer-provided guidelines. Wen et al. introduced a Chained Prior that represents a sequence of task-oriented dialogue guidelines as a Markov chain and proposed GuidedTOD, which updates its policy as the dialogue progresses [16]. By reranking candidate responses with the Chained Prior, GuidedTOD improved guideline compliance over previous methods and increased action-prediction accuracy by approximately 20%. These studies demonstrate the value of expressing desired behavior as explicit guidelines and incorporating them into response selection or generation.

Other studies predict the next response strategy from the user’s state and dialogue history instead of directly applying manually defined guidelines. Cheng et al. proposed MultiESC, which captures fine-grained emotional expressions and their causes and selects a response strategy while anticipating future user reactions [2]. MultiESC uses an A*-inspired lookahead evaluation to select strategies predicted to provide greater long-term support, outperforming previous methods in both strategy planning and response generation. Deng et al. introduced utterance labels and evaluation measures that represent whether the user or the system leads the conversation, and proposed the Knowledge-enhanced Mixed-Initiative model (KEMI), which combines strategy prediction, external-knowledge selection, and response generation [3]. KEMI uses a knowledge graph constructed from mental-health conversations to select strategies and knowledge appropriate to the dialogue context, and it outperformed previous methods in both content preservation and mixed-initiative behavior. Wan et al. proposed EmoDynamix, which models the relationship between fine-grained user emotion states and system response strategies as a heterogeneous graph [15]. EmoDynamix separates strategy prediction from response generation and predicts the next response strategy from relationships in the graph. Experiments on two emotional-support dialogue datasets showed higher strategy-prediction performance than previous methods, reduced bias toward particular strategies, and enabled the strategy-selection process to be traced.

Research on efficient LLM adaptation has also emphasized the quality and utility of training data rather than only its volume. Zhou et al. proposed LIMA, which trains LLaMA on only 1,000 carefully selected instruction–response examples, and showed that even a small amount of high-quality data can teach a model desirable response formats and conversational behavior [21]. In human evaluation, LIMA responses were judged at least as good as GPT-4 responses in 43% of cases and performed strongly against several models trained with much larger alignment datasets. Chen et al. proposed AlpaGasus, which uses a strong LLM to assess each instruction–response example and remove low-quality

data [1]. AlpacaGasus was trained on only 9,000 of Alpaca’s 52,000 examples, yet it outperformed Alpaca trained on the full dataset in both automatic and human evaluation while substantially reducing training time. Liu et al. proposed DEITA, which separately evaluates the complexity, quality, and diversity of instruction data and combines these criteria for selection [8]. Using 6,000 supervised fine-tuning examples—less than one-tenth the amount used by its baselines—DEITA achieved performance comparable to or better than leading open-source models at the time. Xia et al. proposed LESS, which estimates how strongly each training example affects acquisition of a target capability from the gradient similarity between that example and a small set of target-task examples [17]. Training on only 5% of the selection pool could outperform training on the full dataset across several downstream tasks, demonstrating the effectiveness of purpose-driven data selection.

These studies have advanced response-policy control, response-strategy prediction, and efficient selection of training data, but three major challenges remain from the perspective of conversational-style selection. First, guideline-based methods such as DialGuide and GuidedTOD require desirable dialogue policies to be written manually for each purpose or domain. Features expressed implicitly in a corpus—including tone, response timing, and progression across multiple turns—are difficult to specify comprehensively as guidelines, making their design costly. Second, response-strategy prediction methods such as MultiESC, KEMI, and EmoDynamix can choose strategies based on the user state and anticipated future dialogue, but ultimately select a particular strategy at each turn and generate a response from it. Because multiple strategies may be equally valid in the same context, a single prediction cannot fully represent the response diversity of a target corpus. Third, LIMA demonstrates the effectiveness of training on small, high-quality datasets, while AlpacaGasus and DEITA select data according to response quality, complexity, and diversity, and LESS selects data according to their influence on a small set of target-task examples. Even when responses are comparable in general quality or utility for the target task, however, their compatibility with the target style may differ because of their timing or the conversational-state transitions they induce. A selection criterion that does not incorporate dialogue history, response strategy, and state transitions cannot adequately distinguish these differences or reliably preserve the characteristic dialogue flow of the target corpus.

3 BASiS

3.1 Concept

We propose BASiS (Bayesian Style-based Selection), which adapts a local large language model (LLM) to the conversational style required for purpose-specific dialogue through Bayesian data selection based on a small, high-quality corpus. BASiS converts the response policies and dialogue progression expressed implicitly in the small corpus into a form suitable for training a local LLM. Rather than training directly on the small corpus, it extracts conversational states, response strategies, and state transitions, constructs a probabilistic dialogue model, and uses that model as a criterion for selecting data from a large general-purpose corpus.

BASiS evaluates each dialogue instance in the large general-purpose corpus according to its compatibility with the conversational style characteristic of the target corpus. It incorporates the dialogue flow in the small corpus into the training data without requiring detailed, manually written dialogue guidelines for each purpose or domain. By maintaining the conversational state as a probability distribution rather than committing to a single state, BASiS represents settings in which multiple response strategies may be valid for the same context. The selected data are ultimately used for DPO training, which adapts the local LLM to the conversational style of the small corpus.




Fig. 1. Overall architecture of BASiS. The upper pathway extracts conversational features from a small, high-quality corpus and constructs a Bayesian dialogue model. The lower pathway uses the model to score candidate responses from a large general-purpose corpus and adapts a local LLM through LoRA–DPO using the resulting preference data. The numbered circles correspond to the four modules described in the text.

3.2 BASiS Architecture

Figure 1 shows the overall architecture of BASiS. BASiS consists of four modules: the **Dialogue Feature Analysis Module**, **Bayesian Dialogue Modeling Module**, **Style Compatibility Scoring Module**, and **Preference Data Selection and DPO Adaptation Module**.

The upper pathway A in Figure 1, *Style Model Induction*, constructs the style model that serves as the selection criterion. First, the Dialogue Feature Analysis Module (Step 1 in the figure) extracts conversational states, response strategies, and state transitions from dialogue data in the small corpus. The Bayesian Dialogue Modeling Module (Step 2) then aggregates the occurrence patterns of these features and constructs a Bayesian dialogue model consisting of state-transition probabilities and observation likelihoods.

The lower pathway B, *Data Selection & Adaptation*, uses this model to select training data and adapt the LLM. The Style Compatibility Scoring Module (Step 3) receives both the Bayesian dialogue model from the upper pathway and dialogue data from the large general-purpose candidate pool, and assigns each candidate response a compatibility score between 0 and 1. The Preference Data Selection and DPO Adaptation Module (Step 4) uses the scores to construct preference pairs of *chosen* and *rejected* responses, and trains a local LLM with LoRA-based DPO. Thus, the small corpus serves as the style reference, whereas the large corpus serves as the candidate pool for the actual training data. The same pipeline can be applied to other types of purpose-specific dialogue by changing these two corpora.

Dialogue Feature Analysis Module

The Dialogue Feature Analysis Module corresponds to Step 1 in Figure 1. It uses an LLM to analyze each dialogue in the small, high-quality corpus and structures the dialogue progression characteristic of the target corpus. Each turn is represented by three elements: a *conversational state*, a *response strategy*, and a *state transition*.

The conversational state s_t represents the interlocutor’s situation or the stage of the dialogue at turn t . Examples include disclosing emotions or considering solutions in emotional support, attempting a line of reasoning or organizing an error in mathematics tutoring, and presenting symptoms or confirming additional information in medical history-taking. The observation label o_t , which represents the response strategy, describes the policy taken by the respondent in

that state; examples include empathy, questioning, information organization, hints, information provision, and advice. The state transition describes the observed change from s_{t-1} to s_t across a turn in which response strategy o_t is used.

Applying this analysis to the entire small corpus yields a collection of tuples (s_{t-1}, o_t, s_t) . These tuples capture not only surface vocabulary and individual responses, but also the dialogue structure: which strategies tend to be used in particular states and which states tend to follow. This structure converts the conversational style implicit in the corpus into a form that can be used by the subsequent modules.

Bayesian Dialogue Modeling Module

The Bayesian Dialogue Modeling Module constructs a Bayesian dialogue model of the dialogue flow in the target corpus from the features produced by the Dialogue Feature Analysis Module. Let $S = \{s_1, \dots, s_K\}$ denote the set of conversational states and $O = \{o_1, \dots, o_L\}$ the set of observation labels corresponding to response strategies.

The Bayesian dialogue model consists of the state-transition probability $P(s_t | s_{t-1})$ and the observation likelihood $P(o_t | s_t)$. The transition probability represents the probability of moving from conversational state s_{t-1} to s_t in the target corpus. Dialogue flows frequently observed in the target corpus receive high probabilities, whereas rarely observed flows receive low probabilities. The observation likelihood represents the probability that response strategy o_t is observed when the dialogue is in state s_t .

Together, these distributions represent which conversational states are likely to occur in the target corpus and which response strategies are likely to be used in each state. The model does not merely memorize utterances found in the small corpus; instead, it represents probabilistic relationships between states and strategies. It can therefore evaluate the compatibility of new candidate responses that do not appear in the small corpus with the target conversational style.

Style Compatibility Scoring Module

The Style Compatibility Scoring Module uses the Bayesian dialogue model to evaluate a candidate response r_t from the large general-purpose corpus according to its compatibility with the target conversational style. Figure 2 illustrates the five steps: prior-state specification, prediction, response labeling, Bayesian update, and compatibility scoring.

Before a candidate response is presented, the dialogue state is maintained as a probability distribution over the state set S rather than fixed to a single state (Step 1 in Figure 2). In the example, the state distribution through turn $t - 1$, $b_{t-1}(s)$, assigns 0.30 to Emotion disclosure, 0.45 to Emotion processing, and 0.25 to Considering solutions. The predicted state distribution at turn t is then computed using the transition probabilities (Step 2):

$$\hat{b}_t(s_t) = \sum_{s_{t-1} \in S} P(s_t | s_{t-1}) b_{t-1}(s_{t-1}). \quad (1)$$

The state-transition diagram in Step 2 of Figure 2 schematically shows transitions among Disclosure, Processing, and Solutions, together with representative response strategies observed around those transitions. Equation (1) maps the previous distribution b_{t-1} to the predicted distribution $\hat{b}_t$ using $P(s_t | s_{t-1})$, which summarizes the transition patterns in the reference corpus.

Next, an LLM classifies candidate response r_t into observation label o_t using the same label scheme as the Dialogue Feature Analysis Module (Step 3). In the example, the candidate response “That sounds really difficult.” is labeled *Empathy / validation*. The predicted distribution $\hat{b}_t(s_t)$ is then updated according to Bayes’ rule using this observation label and the observation likelihood:

$$b_t(s_t) = \frac{P(o_t | s_t) \hat{b}_t(s_t)}{\sum_{s' \in S} P(o_t | s') \hat{b}_t(s')}. \quad (2)$$




Fig. 2. Schematic example of candidate-response evaluation in the Style Compatibility Scoring Module. The state distribution through the previous turn is predicted using the transition model and then updated with the observation label derived from the candidate response. The blue path indicates where the predicted distribution enters the update equation, and the orange path indicates where the observation label enters it. In this example, the probabilities of the target states Processing and Solutions in the updated distribution are summed to obtain Score = 0.88.

Here, $b_t(s_t)$ is the posterior probability $P(s_t | o_{1:t})$ conditioned on the observation history $o_{1:t}$ after incorporating the observations into the predicted distribution $\hat{b}_t(s_t)$. In Step 4 of Figure 2, the orange path shows that observation o_t determines the likelihood term $P(o_t | s_t)$, whereas the blue path shows where prediction $\hat{b}_t(s_t)$ enters the update. The denominator normalizes the posterior over all states. In this example, the updated distribution assigns 0.12 to Disclosure, 0.66 to Processing, and 0.22 to Solutions.

Finally, let $S_{\text{target}} \subseteq S$ be the set of states defined as desirable for the target conversational style. The compatibility score of candidate response r_t is the total probability assigned to S_{target} in the updated state distribution:

$$\text{Score}(r_t) = \sum_{s \in S_{\text{target}}} b_t(s). \quad (3)$$

The compatibility score ranges from 0 to 1. A larger value indicates that the state distribution after observing the candidate response is more strongly aligned with states considered desirable in the target corpus. In Step 5 of Figure 2, $S_{\text{target}} = \{\text{Processing}, \text{Solutions}\}$, so the score is $0.66 + 0.22 = 0.88$. More generally, S_{target} may include states such as processing emotions or considering appropriate solutions in emotional support, advancing the learner’s own reasoning in mathematics tutoring, or progressively organizing necessary symptom information in medical history-taking.

This module evaluates candidate responses while maintaining a probability distribution over states instead of fixing a single response strategy as the correct next action. Even when each candidate response receives one observation label, responses associated with different strategies can each obtain a high compatibility score if those strategies have high observation likelihoods in the desirable states. Thus, when several response strategies are valid for the same context, the module can evaluate them without restricting the choice to a single strategy.

Preference Data Selection and DPO Adaptation Module

The Preference Data Selection and DPO Adaptation Module selects responses for training from the large general-purpose corpus based on the compatibility scores produced by the Style Compatibility Scoring Module. Among the candidate responses constructed for each dialogue context, a response with high compatibility with the target style is selected as *chosen*, while a response with low compatibility is selected as *rejected*, forming a preference pair.

Table 1. Corpus pairs used for evaluation and the role of each corpus

Dialogue purpose	Small high-quality corpus	Large general corpus
Emotional support	ESConv	DailyDialog
Mathematics tutoring	MathDial	WildChat-1M
Medical history-taking	MediTOD	WildChat-1M

The resulting preference pairs are used to train the local LLM through Direct Preference Optimization (DPO) [12]. DPO directly optimizes the model so that chosen responses are preferred to rejected responses. Rather than updating all model parameters, we apply Low-Rank Adaptation (LoRA) [5] and update only the added low-rank parameters. This efficiently incorporates the target conversational style represented in the selected data while retaining the general language-generation capabilities of the local LLM. Section 4 describes the model, hyperparameters, and amount of training data used in our experiments.

4 Evaluation

4.1 Evaluation Objectives

Our evaluation has two objectives. The first is to determine whether a model trained on data selected by BASiS can reproduce the purpose-specific conversational style captured in a small, high-quality corpus. We compare the model trained on BASiS-selected data against both the untrained base model and a model trained on data randomly sampled from a category coarsely filtered using keywords that describe the target corpus. These comparisons test whether BASiS training strengthens the target style and whether the observed change is associated with BASiS data selection beyond the effects of DPO training and coarse topical matching.

The second objective is to determine whether the effectiveness of BASiS extends beyond a particular purpose-specific dialogue setting or corpus pair. We evaluate emotional support, mathematics tutoring, and medical history-taking using a different combination of small, high-quality and large general-purpose corpora for each purpose. An LLM-based oracle evaluation compares how well each condition reproduces the characteristic conversational style of its target corpus.

4.2 Experimental Method

We use the three corpus pairs shown in Table 1. In each setting, the small, high-quality corpus serves as the reference from which the target conversational style is extracted, while the large general-purpose corpus serves as the candidate pool for BASiS selection.

For emotional support, we use ESConv [7], which contains dialogues between support seekers and supporters, as the small, high-quality corpus, and DailyDialog [6], a collection of everyday multi-turn conversations, as the large general-purpose corpus. For mathematics tutoring, we use MathDial [9], which contains one-to-one mathematics dialogues between teachers and learners simulated by an LLM, as the small corpus, and WildChat-1M [20], which contains real-world interactions between users and LLMs, as the large corpus. For medical history-taking, we use MediTOD [14], which contains clinician-authored history-taking dialogues between doctors and patients, as the small corpus and again use WildChat-1M as the large corpus.

For each setting, we first construct a Bayesian dialogue model from the small, high-quality corpus and use the Style Compatibility Scoring Module to evaluate candidate responses in the large general-purpose corpus. We construct

Table 2. DPO training and training-data selection in each experimental condition

Condition	DPO training	Training-data selection
Base	×	None
BASiS	✓	Compatibility score
Random	✓	Keyword-filtered random

preference data from the resulting compatibility scores and train a local LLM with LoRA-based DPO. To assess the effect of BASiS selection, we also construct the Random condition. For this condition, broad keywords representing the main topics of the small corpus are used to coarsely filter the large corpus by category, after which the training data are sampled randomly from that category. The amount of data and the training procedure match those of BASiS. For each setting, we prepare 100 evaluation prompts and generate one response from each of three models: the untrained base model, the model trained on BASiS-selected data, and the model trained on data randomly sampled from the keyword-filtered category. Using the same prompts across all three conditions provides a paired comparison that controls for differences among prompts.

In the oracle evaluation, an LLM serves as the evaluator and rates how well each model response reproduces the conversational style characteristic of the target corpus. The evaluator is not shown the model names, and the presentation order of the three responses is randomized. We use GPT-5.4 as the oracle LLM with a temperature of 0 to reduce output variability. The oracle LLM receives the evaluation prompt, the preceding dialogue history, and the response to be evaluated. It assigns a score from 1 to 10 on each axis listed in Table 3. The axes are defined with reference to the properties emphasized in each source corpus and the evaluation methods used in related work. They assess not only response content but also purpose-specific response policies, response timing, and dialogue progression across multiple turns. For each axis, we test differences among the Base, BASiS, and Random conditions using the Friedman test. When the omnibus test is significant, we conduct paired permutation tests for BASiS versus Base and BASiS versus Random, with Holm correction for multiple comparisons. We use a significance level of $\alpha = .05$ and report bootstrap 95% confidence intervals to characterize uncertainty in the mean scores.

4.3 Experimental Conditions and Evaluation Measures

We conduct the following two comparisons to distinguish the effect of training with BASiS-selected data from the effect of BASiS data selection itself.

Effect of training with BASiS-selected data: To determine whether training on data selected by BASiS improves reproduction of the target-corpus style, we compare the Base condition, which uses the model before training, against the BASiS condition, which uses the model trained on BASiS-selected data.

Effect of data selection: To determine whether changes in the BASiS condition result from BASiS data selection rather than DPO training or coarse topical matching alone, we compare BASiS against Random. Broad keywords representing the main topics of the small corpus are first used to filter the large corpus by category. Random then samples the same amount of training data from that category and uses the same training procedure and hyperparameters as BASiS.

Table 2 summarizes the three conditions. Base is the model before DPO training. BASiS is trained by DPO on data selected using the compatibility score. Random is trained by DPO on data randomly sampled from a category coarsely identified using broad keywords. All experiments use the same local LLM, Qwen3.5-27B. Base uses the pretrained

Table 3. Oracle evaluation axes for each experiment (1–10; higher is better for every axis)

Evaluation axis	What the axis measures
Corpus pair: ESConv × DailyDialog	
Style Strength	Strength of the ESConv emotional-support style
ESConv Tone Similarity	Similarity to ESConv’s empathetic, accepting, and non-assertive tone
Supporter Role Consistency	Supportive role without lecturing, self-disclosure, or digression
Non-directive Support Style	Respect for the seeker without assertions or imposed directions
Strategy/Stage Alignment	Fit of the response strategy to the current dialogue stage
Premature Advice Avoidance	Advice only after understanding the person’s emotions and situation
Naturalness	Natural, context-appropriate Japanese dialogue
Corpus pair: MathDial × WildChat-1M	
Equitable Tutoring	Space for the learner to think, explain, and solve independently
Learner Reasoning Diagnosis	Recognition of correctness, gaps, and confusion in learner reasoning
Mistake Location and Targeting	Accurate localization of errors or incomplete reasoning steps
Guidance Quality	Accuracy and usefulness of questions, hints, explanations, and examples
Feedback Actionability	Clarity about what the learner should think, calculate, or check next
Answer Revealing Calibration	Appropriate amount and timing of answer revelation
Teacher Move–Stage Alignment	Alignment of the teacher’s response strategy with the dialogue stage
Corpus pair: MediTOD × WildChat-1M	
Response Relevance	Relevance to the patient’s symptoms and preceding context
Overall Quality	Overall relevance, clarity, naturalness, and usefulness
Premature Assessment Avoidance	No diagnosis or treatment planning before sufficient information
Appropriate Uncertainty	Appropriate expression of uncertainty and need for confirmation
Understandability	Clear, patient-friendly questions and explanations
Unsafe Medical Advice Avoidance	Avoidance of potentially harmful medical advice
Unsupported Diagnosis Avoidance	Avoidance of unsupported disease or symptom-cause claims

model without additional training, whereas BASiS and Random use LoRA-based DPO. Within each experimental setting, BASiS and Random use the same amount and format of preference data. Their principal difference is therefore selection by conversational-style compatibility using the Bayesian dialogue model versus random sampling within a coarsely matched topic category. We use a DPO learning rate of 5×10^{-6} and $\beta = 0.1$. LoRA uses rank $r = 8$, $\alpha = 16$, and dropout of 0.05, and training runs for one epoch. These settings are shared across all three corpus pairs.

Table 3 summarizes the seven oracle-evaluation axes used in each experimental setting.

Oracle evaluation axes for ESConv: For ESConv–DailyDialog, the axes capture both the style and progression of emotional-support dialogue. In addition to the strength and tone of ESConv-style support, they assess maintenance of the supporter role, a non-directive response policy, alignment between response strategy and dialogue stage, the timing of advice, and naturalness as Japanese dialogue [3, 7, 10, 19].

Oracle evaluation axes for MathDial: For MathDial–WildChat-1M, the axes evaluate not only mathematical correctness, but also whether a response encourages the learner’s own reasoning and provides pedagogical support appropriate to the learner’s current understanding [9]. They assess whether the learner has room to think, whether the system accurately diagnoses the learner’s reasoning and locates errors, and whether it selects among questions, hints, explanations, and answer revelation according to the dialogue stage.




Fig. 3. Oracle-evaluation results for the ESConv–DailyDialog experiment. Bars show condition means and error bars show bootstrap 95% confidence intervals. Brackets indicate Holm-corrected paired comparisons of BASiS with Base or Random (* $p < .05$, ** $p < .01$, *** $p < .001$).

Oracle evaluation axes for MediTOD: For MediTOD–WildChat-1M, the axes capture the appropriateness of medical history-taking, response quality, and medical safety [14]. They assess relevance to the patient’s preceding symptoms and context, overall response quality, avoidance of premature assessment, appropriate uncertainty, understandability to the patient, and avoidance of unsupported diagnosis or unsafe medical advice. Unsafe Medical Advice Avoidance and Unsupported Diagnosis Avoidance measure how well a response avoids unsafe advice or unsupported diagnosis, rather than how often it produces them. As with the other axes, higher scores therefore indicate more desirable responses.

4.4 Results

This section reports the oracle-evaluation results for the ESConv–DailyDialog, MathDial–WildChat-1M, and MediTOD–WildChat-1M corpus pairs. Figures 3–5 show the mean scores and bootstrap 95% confidence intervals for Base, BASiS, and Random on each axis. Asterisks above the brackets indicate Holm-corrected paired permutation-test differences between BASiS and each comparison condition: * denotes $p < .05$, ** denotes $p < .01$, and *** denotes $p < .001$.

Experimental Setting 1: ESConv × DailyDialog

Figure 3 presents the results for the seven emotional-support axes. BASiS achieved a higher mean score than both Base and Random on every axis. For Style Strength, the mean scores were 8.26 for Base, 8.65 for BASiS, and 8.21 for Random; BASiS was significantly higher than both comparison conditions (both $p < .001$). For ESConv Tone Similarity, the corresponding scores were 8.21, 8.52, and 8.13, and BASiS again significantly outperformed both conditions (both $p < .001$). The oracle evaluation thus identified improvements not only in the overall strength of the emotional-support style but also in its empathetic, accepting, and non-assertive tone.

For Supporter Role Consistency, the scores were 8.34 for Base, 8.60 for BASiS, and 8.41 for Random; BASiS was significantly higher than Base at $p < .01$ and Random at $p < .05$. For Non-directive Support Style, the scores were 7.84, 8.34, and 7.99, respectively; BASiS was significantly higher than Base at $p < .01$ and Random at $p < .05$. These results suggest that, beyond the tone of individual responses, BASiS strengthened the tendency to maintain a supportive role across multiple turns and to continue the conversation without imposing a conclusion on the other person. For




Fig. 4. Oracle-evaluation results for the MathDial–WildChat-1M experiment. Bars show condition means and error bars show bootstrap 95% confidence intervals. Brackets indicate Holm-corrected paired comparisons of BASiS with Base or Random ($***p < .001$).

Strategy/Stage Alignment, the scores were 7.84 for Base, 8.32 for BASiS, and 7.98 for Random. BASiS significantly exceeded Base ($p < .01$), whereas its difference from Random was not significant. For Premature Advice Avoidance, the scores were 8.57, 9.44, and 8.92, respectively; BASiS was significantly higher than Base at $p < .001$ and Random at $p < .05$. Its 0.87-point improvement over Base was the largest among the seven axes. These results suggest a stronger tendency to match response strategies to the current dialogue stage and to offer advice only after understanding the person’s emotions and situation. For Naturalness, the scores were 8.35 for Base, 8.53 for BASiS, and 8.17 for Random. BASiS significantly exceeded Random ($p < .01$), whereas its difference from Base was not significant. The style adaptation achieved by BASiS therefore did not appear to compromise the naturalness of the Japanese dialogue.

Experimental Setting 2: MathDial × WildChat-1M

Figure 4 presents the seven mathematics-tutoring axes. BASiS outperformed both Base and Random on all seven axes, and every paired comparison was significant (all $p < .001$). For Equitable Tutoring, the scores were 6.77 for Base, 7.75 for BASiS, and 6.00 for Random; BASiS exceeded Base by 0.98 points and Random by 1.75 points. The oracle evaluator found that BASiS produced more responses that avoided immediately revealing the answer and instead left room for learners to think, explain, and try a solution. For Learner Reasoning Diagnosis, the scores were 7.38, 8.70, and 6.97, so BASiS exceeded Base by 1.32 points and Random by 1.73 points. For Mistake Location and Targeting, the scores were 7.53, 8.85, and 7.27; BASiS exceeded Base by 1.32 points and Random by 1.58 points. The oracle evaluation identified improvements not only in recognizing whether a learner’s reasoning was correct, but also in locating specific errors or incomplete steps. For Guidance Quality, the scores were 6.88 for Base, 8.02 for BASiS, and 6.57 for Random. For Feedback Actionability, they were 7.01, 8.03, and 6.18, respectively. The 1.85-point gap between BASiS and Random on Feedback Actionability was the largest among the seven axes. The oracle evaluator found that BASiS produced more pedagogical responses that clearly indicated what the learner should think about, calculate, or check next. For Answer Revealing Calibration, the scores were 7.88 for Base, 8.86 for BASiS, and 7.50 for Random. For Teacher Move–Stage Alignment, they were 7.44, 8.62, and 7.06, respectively. These results suggest a stronger tendency to align the amount and timing of answer revelation, as well as teacher moves such as questions, focusing, and explanations, with the learner’s state and the dialogue stage.




Fig. 5. Oracle-evaluation results for the MediTOD–WildChat-1M experiment. Bars show condition means and error bars show bootstrap 95% confidence intervals. Brackets indicate Holm-corrected paired comparisons of BASiS with Base or Random (* $p < .05$, ** $p < .01$, *** $p < .001$).

Experimental Setting 3: MediTOD $\times$ WildChat-1M

Figure 5 presents the seven medical history-taking axes. BASiS had the highest mean score on all seven axes, although the pattern of statistical significance differed across axes. For Response Relevance, the scores were 4.46 for Base, 5.09 for BASiS, and 4.91 for Random. For Overall Quality, the corresponding scores were 5.16, 5.69, and 5.46. BASiS significantly exceeded Base on both axes (both $p < .01$), whereas neither difference from Random was significant. Because Random samples from a coarsely matched topic category, it may already preserve some response relevance and overall quality; these two axes do not establish an advantage for BASiS over Random. For Premature Assessment Avoidance, the scores were 8.43 for Base, 8.90 for BASiS, and 8.63 for Random. BASiS had the highest mean, but neither Holm-corrected paired comparison was significant. For Appropriate Uncertainty, the scores were 7.84, 8.26, and 7.92, respectively, and BASiS significantly exceeded both Base and Random (both $p < .01$). Thus, the statistically supported finding is that BASiS more consistently expressed uncertainty and the need for further confirmation when the available information was insufficient. For Understandability, the scores were 8.56 for Base, 8.77 for BASiS, and 8.55 for Random; BASiS significantly exceeded both conditions (both $p < .05$). For Unsafe Medical Advice Avoidance, the scores were 9.52, 9.71, and 9.36, respectively; BASiS was significantly higher than Base at $p < .001$ and Random at $p < .05$. For Unsupported Diagnosis Avoidance, the scores were 9.62, 9.82, and 9.61, and BASiS significantly exceeded both conditions (both $p < .05$). These results suggest that BASiS strengthened response policies that remain understandable to patients while avoiding poorly grounded medical advice and diagnosis.

4.5 Discussion

We evaluated BASiS on three distinct types of purpose-specific dialogue: emotional support, mathematics tutoring, and medical history-taking. In the oracle evaluation, BASiS achieved a higher mean score than Base and Random on all 21 axes across the three settings. The pattern of statistical significance differed by domain and axis: improvements were broad in emotional support and mathematics tutoring, whereas response relevance, overall quality, and safety showed different patterns in medical history-taking. We discuss the implications of these results below.

Discussion 1: Reproducing conversational style in emotional-support dialogue

In the ESConv–DailyDialog experiment, BASiS achieved the highest mean on all seven axes. The largest improvement occurred in Premature Advice Avoidance, which measures whether advice is offered only after the person’s emotions and situation have been understood. Non-directive Support Style measures whether the response respects the support seeker’s agency without imposing assertions or directions. BASiS also significantly exceeded both comparison conditions on Style Strength and ESConv Tone Similarity. The oracle results therefore indicate changes in both the empathetic language characteristic of ESConv and the *timing* of advice. BASiS significantly exceeded both comparison conditions on Supporter Role Consistency and exceeded Base on Strategy/Stage Alignment. Its effect may therefore extend beyond surface expressions in a single turn to response policies that follow dialogue progression. On Naturalness, BASiS significantly exceeded Random and did not differ significantly from Base, providing no evidence that style adaptation compromised the naturalness of the Japanese dialogue. These findings are consistent with the BASiS design, which represents conversational style through relationships between dialogue states and response strategies rather than vocabulary alone.

Discussion 2: Reproducing pedagogical dialogue policies in mathematics tutoring

In the MathDial–WildChat-1M experiment, BASiS significantly outperformed both Base and Random on all seven axes. Effective mathematics tutoring requires more than producing the correct answer: a system must understand whether the learner’s reasoning is correct, incomplete, or confused (Learner Reasoning Diagnosis) and identify the location of errors or incomplete steps (Mistake Location and Targeting). It must also calibrate the amount and timing of answer revelation to the learner’s understanding (Answer Revealing Calibration) and align response strategies such as questions, hints, and explanations with the current dialogue stage (Teacher Move–Stage Alignment). Improvements across all axes, including these measures, suggest that the pedagogical dialogue progression in the target corpus was reflected in the BASiS condition. The largest difference between BASiS and Random occurred in Feedback Actionability, which measures whether the learner is given a clear next action to think through, calculate, or verify. This result suggests that random sampling within a coarsely matched topic category may retain topical relevance but insufficient pedagogical structure to guide the learner’s next action. By selecting data compatible with the transitions between conversational states and response strategies extracted from MathDial, BASiS appears to have incorporated responses that lead into the learner’s next reasoning or calculation step. Random scored below Base on all seven axes. Because we did not directly test the difference between Base and Random, however, we cannot conclude that random-data DPO reduced pedagogical response quality. The fact that BASiS outperformed both Base and Random nevertheless indicates that how the training data are selected matters in addition to whether DPO is performed.

Discussion 3: Response quality and safety in medical history-taking

In the MediTOD–WildChat-1M experiment, BASiS achieved the highest mean on all seven axes. It significantly exceeded both Base and Random on Appropriate Uncertainty, which measures whether uncertainty and the need for confirmation are expressed appropriately; Understandability, which measures clarity to the patient; Unsafe Medical Advice Avoidance; and Unsupported Diagnosis Avoidance. These results suggest that BASiS strengthened safety-conscious response policies without sacrificing patient-friendly language. On Response Relevance and Overall Quality, BASiS significantly exceeded Base but not Random. Because Random samples from a broad keyword-filtered category, topical matching may already preserve some relevance and overall quality. BASiS also had the highest mean on Premature Assessment Avoidance, but neither paired difference was significant. Its advantage therefore cannot be assumed to be equally strong across all aspects of medical history-taking.

Discussion 4: Selecting dialogue progression rather than response content alone

Across the three settings, BASiS improved not only general response quality but also axes associated with response timing and alignment with the dialogue stage. In emotional support, it exceeded both comparison conditions in the appropriateness of moving to advice only after understanding the person’s emotions and situation (Premature Advice Avoidance), and exceeded Base in alignment between response strategy and dialogue stage (Strategy/Stage Alignment). In mathematics tutoring, it exceeded both comparison conditions in the alignment of answer-revelation timing with learner understanding (Answer Revealing Calibration) and in alignment between teacher moves and dialogue stage (Teacher Move–Stage Alignment). In medical history-taking, it exceeded both comparison conditions in appropriate expression of uncertainty and the need for further confirmation (Appropriate Uncertainty) and in avoiding diagnosis when information was insufficient (Unsupported Diagnosis Avoidance). None of these axes can be determined solely from the content of an isolated response; they require the preceding history and current dialogue stage to be considered. Naturalness in emotional support and Response Relevance, Overall Quality, and Understandability in medical history-taking also improved over at least one comparison condition. BASiS may therefore strengthen purpose-specific response policies while maintaining or improving general naturalness and understandability. This result is consistent with the BASiS design in Equation (3), which evaluates candidate responses using posterior probabilities based on state transitions and observation labels rather than general single-turn quality. Using transitions between conversational states and response strategies extracted from a small, high-quality corpus as the selection criterion may have allowed BASiS to capture the target corpus’s characteristic pattern of *when* and *how* to respond.

Discussion 5: Contribution of BASiS data selection

BASiS achieved a higher mean score than Random on every axis in all three settings. The BASiS and Random conditions use the same base model, amount of training data, DPO procedure, and LoRA configuration. Random is not sampled without constraint from the full large corpus; it is sampled from a category first identified using broad keywords that describe the small corpus. The comparison therefore tests the additional contribution of using compatibility based on conversational states, response strategies, and state transitions beyond coarse topical matching. In the MathDial setting in particular, Random scored below Base on every axis, whereas BASiS scored above Base on every axis. This contrast suggests that, even among data that broadly match the target topic, the choice of training examples can change how well the model adapts to a purpose-specific conversational style. When a large general-purpose corpus is used for purpose-specific dialogue, it is therefore important not only to filter by topic, but also to select data aligned with the dialogue progression of the target corpus.

Discussion 6: Applicability to three types of purpose-specific dialogue

By changing the small, high-quality reference corpus and the large general-purpose selection corpus, we applied the same BASiS pipeline to emotional support, mathematics tutoring, and medical history-taking. BASiS achieved the highest mean on every axis in all three domains, suggesting that the method is not limited to the emotional-support style of ESConv and may be able to extract styles from different target corpora and use them to select training data. At the same time, some emotional-support and medical history-taking axes showed no significant difference from Random, indicating that the effect of BASiS does not emerge equally strongly in every domain or for every component of conversational style. The aspects that are most strongly reflected may depend on which states are defined as the desirable set S_{target} and which response strategies are extracted as observation labels for each target corpus. Thus, although the overall BASiS pipeline can be shared across domains, future work must strengthen how states, strategies, and the desirable state set are defined from the characteristics of each purpose-specific dialogue setting.

5 Conclusion

We proposed BASiS (Bayesian Style-based Selection), a method that adapts a local large language model (LLM) to the conversational style required for purpose-specific dialogue through Bayesian data selection based on a small, high-quality corpus. BASiS extracts conversational states, response strategies, and state transitions from the small corpus and represents them as a Bayesian dialogue model. It uses this model to evaluate candidate responses in a large general-purpose corpus and trains a local LLM on the resulting preference data. This mechanism uses the response policies and multi-turn dialogue progression captured in a small corpus as criteria for selecting training data, without requiring detailed dialogue guidelines to be written manually for each purpose. By representing conversational states as probability distributions, it can select data suited to the target purpose while preserving the diversity in which multiple response strategies may be valid for the same context.

In comparative experiments on three distinct types of purpose-specific dialogue, BASiS achieved the highest mean score on every evaluation axis. The results suggest that BASiS can incorporate into a model, according to the purpose, empathetic responses and appropriate advice timing, pedagogical support adapted to the learner’s understanding, and understandable response policies that avoid unwarranted certainty. Comparison with a model trained on data randomly sampled from a coarsely keyword-filtered category further suggests the importance of selecting training data according to the target conversational style rather than relying only on broad topical matching.

Future work will evaluate applicability to other models and investigate how to design conversational states, response strategies, and desirable state sets for different purposes. Through these efforts, we aim to extend BASiS to a broad range of purpose-specific dialogue applications, including counseling, education, and medicine.

References

- [1] Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, and Hongxia Jin. 2024. AlpaGasus: Training a Better Alpaca with Fewer Data. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=FdVXgSJhvz>
- [2] Yi Cheng, Wenge Liu, Wenjie Li, Jiashuo Wang, Ruihui Zhao, Bang Liu, Xiaodan Liang, and Yefeng Zheng. 2022. Improving Multi-turn Emotional Support Dialogue Generation with Lookahead Strategy Planning. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 3014–3026. doi:10.18653/v1/2022.emnlp-main.195
- [3] Yang Deng, Wenxuan Zhang, Yifei Yuan, and Wai Lam. 2023. Knowledge-enhanced Mixed-initiative Dialogue System for Emotional Support Conversations. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Toronto, Canada, 4079–4095. doi:10.18653/v1/2023.acl-long.225
- [4] Prakhar Gupta, Yang Liu, Di Jin, Behnam Hedayatnia, Spandana Gella, Sijia Liu, Patrick Lange, Julia Hirschberg, and Dilek Hakkani-Tur. 2023. DialGuide: Aligning Dialogue Model Behavior with Developer Guidelines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 14031–14047. doi:10.18653/v1/2023.findings-emnlp.935
- [5] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. The Tenth International Conference on Learning Representations (ICLR 2022). <https://openreview.net/forum?id=nZeVKeeFYf9>
- [6] Yanran Li, Hui Su, Xiaoyu Shen, Wenjie Li, Ziqiang Cao, and Shuzi Niu. 2017. DailyDialog: A Manually Labelled Multi-turn Dialogue Dataset. In *Proceedings of the Eighth International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Asian Federation of Natural Language Processing, Taipei, Taiwan, 986–995. <https://aclanthology.org/I17-1099/>
- [7] Siyang Liu, Chujie Zheng, Orianna Demasi, Sahand Sabour, Yu Li, Zhou Yu, Yong Jiang, and Minlie Huang. 2021. Towards Emotional Support Dialog Systems. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Association for Computational Linguistics, Online, 3469–3483. doi:10.18653/v1/2021.acl-long.269
- [8] Wei Liu, Weihao Zeng, Keqing He, Yong Jiang, and Junxian He. 2024. What Makes Good Data for Alignment? A Comprehensive Study of Automatic Data Selection in Instruction Tuning. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=BTKAeLqLMw>
- [9] Jakub Macina, Nico Daheim, Sankalan Pal Chowdhury, Tanmay Sinha, Manu Kapur, Iryna Gurevych, and Mrinmaya Sachan. 2023. MathDial: A Dialogue Tutoring Dataset with Rich Pedagogical Properties Grounded in Math Reasoning Problems. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 5602–5621. doi:10.18653/v1/2023.findings-emnlp.372

- [10] Remi Mir, Bjarke Felbo, Nick Obradovich, and Iyad Rahwan. 2019. Evaluating Style Transfer for Text. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, 495–504. doi:10.18653/v1/N19-1049
- [11] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems*, Vol. 35. Curran Associates, Inc., Red Hook, NY, USA, 27730–27744.
- [12] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 53728–53741. doi:10.52202/075280-2338
- [13] Hannah Rashkin, Eric Michael Smith, Margaret Li, and Y-Lan Boureau. 2019. Towards Empathetic Open-domain Conversation Models: A New Benchmark and Dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, Florence, Italy, 5370–5381. doi:10.18653/v1/P19-1534
- [14] Vishal Vivek Saley, Goonjan Saha, Rocktim Jyoti Das, Dinesh Raghu, and Mausam. 2024. MediTOD: An English Dialogue Dataset for Medical History Taking with Comprehensive Annotations. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Miami, Florida, USA, 16843–16877. doi:10.18653/v1/2024.emnlp-main.936
- [15] Chenwei Wan, Matthieu Labeau, and Chloé Clavel. 2025. EmoDynamiX: Emotional Support Dialogue Strategy Prediction by Modelling MiXed Emotions and Discourse Dynamics. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics, Albuquerque, New Mexico, 1678–1695. doi:10.18653/v1/2025.naacl-long.81
- [16] Xiangyu Wen, Jianyuan Zhong, Zhijian Xu, and Qiang Xu. 2025. Guideline Compliance in Task-Oriented Dialogue: The Chained Prior Approach. In *Findings of the Association for Computational Linguistics: NAACL 2025*. Association for Computational Linguistics, Albuquerque, New Mexico, 6765–6791. doi:10.18653/v1/2025.findings-naacl.377
- [17] Mengzhou Xia, Sadhika Malladi, Suchin Gururangan, Sanjeev Arora, and Danqi Chen. 2024. LESS: Selecting Influential Data for Targeted Instruction Tuning. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, Vienna, Austria, 54104–54132. <https://proceedings.mlr.press/v235/xia24c.html>
- [18] Bolin Zhang, Jiahao Wang, Qianlong Du, Jiajun Zhang, Zhiying Tu, and Dianhui Chu. 2025. A Survey on Data Selection for LLM Instruction Tuning. *Journal of Artificial Intelligence Research* 83, Article 32 (2025), 14 pages. doi:10.1613/jair.1.17625
- [19] Saizheng Zhang, Emily Dinan, Jack Urbanek, Arthur Szlam, Douwe Kiela, and Jason Weston. 2018. Personalizing Dialogue Agents: I have a dog, do you have pets too?. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Melbourne, Australia, 2204–2213. doi:10.18653/v1/P18-1205
- [20] Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. 2024. WildChat: 1M ChatGPT Interaction Logs in the Wild. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=Bl8u7ZRlbM>
- [21] Chunting Zhou, Pengfei Liu, Puxin Xu, Srinu Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, Lili Yu, Susan Zhang, Gargi Ghosh, Mike Lewis, Luke Zettlemoyer, and Omer Levy. 2023. LIMA: Less Is More for Alignment. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 55006–55021. doi:10.52202/075280-2400