

BASiS: Transferring Dialogue Progression to Language Models through Bayesian Data Selection

ANONYMOUS AUTHOR(S)

We propose BASiS (Bayesian Style-based Selection), a framework for scaling the interactional knowledge encoded in a small, high-quality dialogue corpus to large-scale LLM training. We focus on settings where dialogue must progress according to the user’s state and the stage of the interaction as exemplified in emotional support, education, and healthcare. In such settings, dialogue quality depends not only on utterance content but also on how the dialogue develops across turns, which involves procedural interactional knowledge of how to adapt response strategies as the interaction evolves. Although such knowledge is implicitly encoded in small, high-quality corpora, these corpora alone provide limited data for LLM training. To overcome this challenge, BASiS models the progression patterns in the high-quality corpus and uses the resulting model to guide training-data selection from a large general-purpose corpus. It extracts dialogue states, response strategies, and state transitions from the small corpus and represents the probabilistic relationships as a Bayesian dialogue model. The model scores candidate responses in a large corpus according to how well they support the dialogue progression exhibited in the high-quality corpus, and the scores are used to construct preference pairs for LLM fine-tuning with Direct Preference Optimization (DPO). We evaluated BASiS in emotional support, mathematics tutoring, and medical history-taking against the base model and a model trained on randomly selected in-category data. In the human evaluation, BASiS achieved the highest mean score on all 19 evaluation dimensions and was selected most often in all three domains. In the LLM-based oracle evaluation, BASiS achieved the highest mean score on all 21 evaluation dimensions, with broad statistically significant improvements in emotional support and mathematics tutoring.

CCS Concepts: • **Computing methodologies** → **Natural language generation**; **Discourse, dialogue and pragmatics**; Learning from implicit feedback; • **Human-centered computing** → *Natural language interfaces*.

Additional Key Words and Phrases: dialogue progression, data selection, Bayesian dialogue modeling, response strategy, preference optimization, emotional support dialogue, large language models

ACM Reference Format:

Anonymous Author(s). 2026. BASiS: Transferring Dialogue Progression to Language Models through Bayesian Data Selection. In *Proceedings of 32nd Annual ACM Conference on Intelligent User Interfaces (IUI '27)*. ACM, New York, NY, USA, 17 pages.

1 Introduction

Advances in instruction tuning and reinforcement learning from human feedback (RLHF) have enabled recent large language models (LLMs) to follow instructions, generate accurate responses, and sustain natural and coherent multi-turn conversations across a wide range of topics [10]. In settings such as emotional support, mathematics tutoring, and medical history-taking, however, a system must do more than provide correct information: it must select a response strategy according to the interlocutor’s state and the current stage of the interaction [6, 12]. For example, it may need to acknowledge a person’s feelings before offering advice, encourage a learner to reason independently, or gather

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

sufficient symptoms before drawing a conclusion. Effective responses in these settings therefore require the system to control dialogue progression across turns.

LLM-based response generation has approached this requirement in two main ways. Guideline-based methods specify desirable response policies in natural language and guide generation to follow them [3, 16]. Data-driven methods instead adapt model behavior through instruction tuning or by selecting high-quality training examples [18, 21]. The former requires dialogue policies to be designed manually for each setting, making it difficult to enumerate implicit features such as response timing and multi-turn progression. The latter depends on sufficient training data, whereas detailed interactional knowledge is often encoded in small, carefully constructed dialogue corpora that alone provide too few examples for LLM training.

A separate line of research has modeled dialogue progression through dialogue actions, response strategies, and their transitions. Ganesh et al. predicted the next teacher talk move in classroom dialogue [2]; Wan et al. modeled relationships between user emotion states and response strategies as a heterogeneous graph [15]; and Rudolph et al. regularized next-dialogue-act prediction using a transition matrix derived from counseling conversations [13]. These studies capture important structure beyond surface utterances, but primarily evaluate action or strategy prediction. Their integration with LLMs for natural-language response generation, and evaluation of whether the modeled progression appears in generated responses, remain limited.

To bridge these two lines of work, we propose BASiS (Bayesian Style-based Selection), a method for transferring dialogue progression from a small, high-quality corpus to a local LLM. BASiS first induces a progression model comprising conversational states, response strategies, state transitions, and their probabilistic relationships. It then scores candidate responses in a large general-purpose corpus according to their compatibility with that progression and constructs preference pairs from the scores. Finally, it adapts the local LLM using Direct Preference Optimization (DPO) [11] with Low-Rank Adaptation (LoRA) [4]. This combination uses large-corpus selection to scale the interactional knowledge encoded in the small corpus, while reducing the need to manually author detailed dialogue guidelines. By maintaining a probability distribution over conversational states, BASiS can account for dialogue history and multiple valid response strategies. We evaluate the same pipeline in emotional support, mathematics tutoring, and medical history-taking by changing the reference corpus and candidate pool.

The remainder of this paper is organized as follows. Section 2 reviews related work on guideline-based response control, dialogue-history-based prediction of dialogue actions and response strategies, and training-data selection for LLMs. Section 3 describes the BASiS training pipeline and the role of each module. Section 4 presents the experiments used to evaluate BASiS, their results, and our discussion. Finally, Section 5 concludes the paper.

2 Related Work

One approach to controlling behavior in dialogue systems is to provide desirable response policies as natural-language guidelines and control response generation accordingly. Gupta et al. proposed DialGuide, which describes in natural language both the conversational contexts in which a guideline applies and the content that a response should include [3]. DialGuide addresses guideline selection, guideline-grounded response generation, and verification that generated responses satisfy the selected guideline. Experiments on open-domain and safety-related dialogue data showed that it can generate safe and engaging responses that follow developer-provided guidelines. Wen et al. introduced a Chained Prior that represents a sequence of dialogue guidelines as a Markov chain and proposed GuidedTOD, which updates its policy as the dialogue progresses [16]. By reranking candidate responses with the Chained Prior, GuidedTOD improved guideline compliance over previous methods and increased action-prediction accuracy by approximately 20%. These

studies demonstrate the value of expressing desired behavior as explicit guidelines and incorporating them into response selection or generation.

Other studies predict the next dialogue action or response strategy from the preceding interaction instead of directly applying manually defined guidelines. Ganesh et al. introduced Future Talk Move Prediction, which predicts a teacher’s next talk move from the classroom dialogue history and the preceding talk moves [2]. By formulating teacher behavior as multiclass classification, their model substantially outperformed multiple baselines and exhibited patterns similar to those observed in human predictions. Wan et al. proposed EmoDynamix, which models discourse-level relationships between fine-grained user emotion states and system response strategies as a heterogeneous graph [15]. Across two emotional-support datasets, it achieved higher F1 scores and lower strategy-preference bias than previous methods, while also supporting interpretable backtracing of its predictions. Rudolph et al. incorporated a transition matrix estimated from counseling dialogues as a KL regularizer for next-dialogue-act prediction, aligning the predicted distribution with observed dialogue progression [13]. Their method improved macro-F1 by 9–42% relative across multiple encoders and also improved progression alignment.

Research on efficient LLM adaptation has also emphasized the quality and task relevance of training data rather than only its volume. LIMA showed that a small set of high-quality instruction–response examples can teach desirable response formats and conversational behavior [21]. AlpaGasus and DEITA select data according to criteria such as quality, complexity, and diversity, achieving strong performance with less training data [1, 7]. LESS further demonstrated the effectiveness of selecting examples according to their estimated influence on target-task performance [17]. Together, these findings show that selecting high-quality data suited to the target is important for efficient model adaptation.

These studies have advanced response-policy control, dialogue-action and response-strategy prediction, and efficient selection of training data. Nevertheless, two challenges remain when incorporating corpus-derived dialogue progression into actual response generation. First, guideline-based methods such as DialGuide and GuidedTOD require desirable dialogue policies to be designed manually whenever the dialogue setting changes. Features expressed implicitly in a corpus—including tone, response timing, policy changes based on the interlocutor’s state, and progression across multiple turns—are difficult to specify comprehensively as guidelines. This imposes substantial design costs. Moreover, because omitted features are unlikely to affect response control, the scope and quality of that control depend on the completeness of the guidelines. Second, existing work that models dialogue actions or response strategies primarily evaluates prediction rather than natural-language generation based on the predicted structure. Ganesh et al. output the label of the teacher’s next talk move but do not generate a teacher response that realizes that move. EmoDynamix can pass its predicted strategy to an external generator, but its reported evaluation concerns strategy-prediction F1, strategy-preference bias, and interpretability rather than the quality of generated emotional-support responses. Rudolph et al. likewise evaluate next-dialogue-act prediction and alignment with dialogue progression, explicitly leaving end-to-end integration with response generation unexplored. Interactional knowledge learned from a corpus must therefore be incorporated into a response generator and evaluated at the level of the generated responses.

3 BASiS

3.1 Concept

We propose BASiS (Bayesian Style-based Selection), which transfers the dialogue progression encoded in a small, high-quality corpus to a local large language model (LLM) through Bayesian data selection. BASiS converts the response policies and progression patterns expressed implicitly in the small corpus into a form suitable for training the local

LLM. Rather than training directly on the small corpus, it extracts conversational states, response strategies, and state transitions, constructs a probabilistic dialogue model, and uses that model as a criterion for selecting data from a large general-purpose corpus. BASiS evaluates each dialogue instance in the large general-purpose corpus according to its progression compatibility with the target corpus. It incorporates the interactional knowledge in the small corpus into the training data without requiring detailed, manually written dialogue guidelines for each setting. By maintaining the conversational state as a probability distribution rather than committing to a single state, BASiS represents settings in which multiple response strategies may be valid for the same context. The selected data are ultimately used for DPO training, which transfers the progression patterns in the small corpus to the local LLM.

3.2 BASiS Training Pipeline

Figure 1 shows the BASiS training pipeline. The pipeline comprises four modules—the **Dialogue Feature Analysis Module**, **Bayesian Dialogue Modeling Module**, **Dialogue Progression Scoring Module**, and **Preference Data Selection and DPO Adaptation Module**—and covers the process from progression-model induction through training-data selection to DPO-based model adaptation. The upper pathway A in Figure 1, *Dialogue Progression Model Induction*, constructs the progression model used as the selection criterion from the small, high-quality corpus. First, the Dialogue Feature Analysis Module (Step 1) extracts conversational states, response strategies, and state transitions as conversational features from the dialogue data in the small corpus. The Bayesian Dialogue Modeling Module (Step 2) then aggregates the occurrence patterns of these features and constructs a Bayesian dialogue model consisting of state-transition probabilities and observation likelihoods. The lower pathway B, *Data Selection & Adaptation*, uses this model to select training data and adapt the LLM. The Dialogue Progression Scoring Module (Step 3) receives both the Bayesian dialogue model from the upper pathway and dialogue data from the large general-purpose candidate pool, and assigns each candidate response a progression-compatibility score between 0 and 1. The Preference Data Selection and DPO Adaptation Module (Step 4) uses the scores to construct preference pairs of *chosen* and *rejected* responses, and trains a local LLM with LoRA-based DPO. Thus, the small corpus serves as the progression reference, whereas the large corpus serves as the candidate pool for the actual training data. The same pipeline can be applied in other dialogue settings by changing these two corpora.

Dialogue Feature Analysis Module

The Dialogue Feature Analysis Module corresponds to Step 1 in Figure 1. It uses an LLM to analyze the small, high-quality corpus and generate a label scheme that represents its dialogue progression. We use GPT-5.4-pro with the API-default temperature and high reasoning effort. The model receives the dialogue text and existing annotations.

The prompt requests 4–7 latent *conversational states*, *response strategies* observed in each state, desirable and off-purpose state groups, initial-state probabilities, state-transition probabilities, and observation probabilities. States describe the interlocutor’s situation or dialogue stage, whereas strategies describe the respondent’s observable policy; the prompt therefore requires the two to remain conceptually distinct. This is constrained ontology induction: rather than fixing the label set manually or allowing the LLM to generate labels without constraints, humans specify the number and roles of labels and the output format. For example, when applied to ESConv, which we use in our experiments, the analysis produced *opening_rapport*, a state that creates a safe atmosphere through greetings and brief check-ins, and *empathic_support*, a state that acknowledges the help seeker’s emotions to alleviate isolation and psychological burden. It also produced *explore_question*, a response strategy that asks about events and background in context to understand the seeker’s situation and concerns, and *reflect_validate*, a strategy that paraphrases and validates the seeker’s words and emotions to convey empathy and acceptance.




Fig. 1. BASiS training pipeline. The upper pathway extracts conversational features from a small, high-quality corpus and constructs a Bayesian dialogue model. The lower pathway uses the model to score candidate responses from a large general-purpose corpus and adapts a local LLM through LoRA–DPO using the resulting preference data. The numbered circles correspond to the four modules described in the text.

Bayesian Dialogue Modeling Module

The Bayesian Dialogue Modeling Module uses the conversational states, response strategies, and state transitions produced by the Dialogue Feature Analysis Module to construct a Bayesian dialogue model representing the dialogue progression of the target corpus. Let $S = \{s_1, \dots, s_K\}$ denote the conversational states and $O = \{o_1, \dots, o_L\}$ the observation labels corresponding to response strategies. The model consists of the initial-state probability $P(s_1)$, state-transition probability $P(s_t | s_{t-1})$, and observation likelihood $P(o_t | s_t)$. The transition probabilities represent the relative tendency, estimated from the target corpus, to move from s_{t-1} to s_t ; the observation likelihood represents the probability of observing strategy o_t in state s_t . These probabilities are normalized from values estimated by the LLM as overall tendencies in the small corpus. Together, the probability distributions represent which dialogue states are likely to occur and which response strategies are likely to be used in each state. Because the resulting model represents probabilistic relationships between states and strategies rather than memorizing utterances in the small corpus, it can also be applied to candidate responses not contained in that corpus.

Dialogue Progression Scoring Module

The Dialogue Progression Scoring Module uses the Bayesian dialogue model to evaluate a candidate response r_t from the large general-purpose corpus according to its compatibility with the dialogue progression in the target corpus. Figure 2 illustrates the five steps: prior-state specification, prediction, response labeling, Bayesian update, and progression-compatibility scoring.

Before a candidate response is presented, the dialogue state is maintained as a probability distribution over the state set S rather than fixed to a single state (Step 1 in Figure 2). In the example, the state distribution through turn $t - 1$, $b_{t-1}(s)$, assigns 0.30 to Emotion disclosure, 0.45 to Emotion processing, and 0.25 to Considering solutions. The predicted state distribution at turn t is then computed using the transition probabilities (Step 2):

$$\hat{b}_t(s_t) = \sum_{s_{t-1} \in S} P(s_t | s_{t-1}) b_{t-1}(s_{t-1}). \quad (1)$$




Fig. 2. Schematic example of candidate-response evaluation in the Dialogue Progression Scoring Module. The state distribution through the previous turn is predicted using the transition model and then updated with the observation label derived from the candidate response. The blue path indicates where the predicted distribution enters the update equation, and the orange path indicates where the observation label enters it. In this example, the probabilities of the target states Processing and Solutions in the updated distribution are summed to obtain Score = 0.88.

The state-transition diagram in Step 2 of Figure 2 schematically shows transitions among Disclosure, Processing, and Solutions, together with representative response strategies observed around those transitions. Equation (1) maps the previous distribution b_{t-1} to the predicted distribution $\hat{b}_t$ using $P(s_t | s_{t-1})$, which summarizes the transition patterns in the reference corpus.

Next, GPT-5.4 classifies candidate response r_t into an observation label o_t from the fixed scheme (Step 3). We use the API-default temperature and high reasoning effort. The prompt provides the dialogue history, candidate response, and allowed strategy labels with their descriptions; it requires exactly one closest label and forbids new labels or state names. The model returns a JSON object containing the observation label, a confidence value from 0 to 1, and a brief rationale. This confidence is auxiliary classification information and is not the progression-compatibility score in Equation (3). In the example, “That sounds really difficult.” is labeled *Empathy / validation*. The LLM does not assign a latent state directly; the predicted state distribution is updated by Bayes’ rule using the observation label and likelihood:

$$b_t(s_t) = \frac{P(o_t | s_t) \hat{b}_t(s_t)}{\sum_{s' \in S} P(o_t | s') \hat{b}_t(s')}. \quad (2)$$

Here, $b_t(s_t)$ is the posterior probability $P(s_t | o_{1:t})$ conditioned on the observation history $o_{1:t}$ after incorporating the observations into the predicted distribution $\hat{b}_t(s_t)$. In Step 4 of Figure 2, the orange path shows that observation o_t determines the likelihood term $P(o_t | s_t)$, whereas the blue path shows where prediction $\hat{b}_t(s_t)$ enters the update. The denominator normalizes the posterior over all states. In this example, the updated distribution assigns 0.12 to Disclosure, 0.66 to Processing, and 0.22 to Solutions.

Finally, let $S_{\text{target}} \subseteq S$ be the set of states defined as desirable for the target dialogue progression. The progression-compatibility score of candidate response r_t is the total probability assigned to S_{target} in the updated state distribution:

$$\text{Score}(r_t) = \sum_{s \in S_{\text{target}}} b_t(s). \quad (3)$$

The progression-compatibility score ranges from 0 to 1. A larger value indicates that the state distribution after observing the candidate response is more strongly aligned with states considered desirable in the target corpus. In Step 5 of Figure 2, $S_{\text{target}} = \{\text{Processing, Solutions}\}$, so the score is $0.66 + 0.22 = 0.88$. More generally, S_{target} may include states such as processing emotions or considering appropriate solutions in emotional support, advancing the learner’s own reasoning in mathematics tutoring, or progressively organizing necessary symptom information in medical history-taking.

This module evaluates candidate responses while maintaining a probability distribution over states instead of fixing a single response strategy as the correct next action. Even when each candidate response receives one observation label, responses associated with different strategies can each obtain a high progression-compatibility score if those strategies have high observation likelihoods in the desirable states. Thus, when several response strategies are valid for the same context, the module can evaluate them without restricting the choice to a single strategy.

Preference Data Selection and DPO Adaptation Module

The Preference Data Selection and DPO Adaptation Module selects responses for training from the large general-purpose corpus based on the progression-compatibility scores produced by the Dialogue Progression Scoring Module. Among the candidate responses constructed for each dialogue context, a response with high compatibility with the target progression is selected as *chosen*, while a response with low compatibility is selected as *rejected*, forming a preference pair. The resulting preference pairs are used to train the local LLM through Direct Preference Optimization (DPO) [11]. DPO directly optimizes the model so that chosen responses are preferred to rejected responses. Rather than updating all model parameters, we apply Low-Rank Adaptation (LoRA) [4] and update only the added low-rank parameters. This efficiently incorporates the dialogue progression represented in the selected data while retaining the general language-generation capabilities of the local LLM.

4 Evaluation

4.1 Evaluation Objectives

Our evaluation has two objectives. The first is to determine whether a model trained on data selected by BASiS can reproduce the dialogue progression captured in a small, high-quality corpus. We compare the model trained on BASiS-selected data against both the untrained base model and a model trained on data randomly sampled from a category coarsely filtered using keywords that describe the target corpus. These comparisons test whether BASiS training strengthens the target response policy and whether the observed change is associated with BASiS data selection beyond the effects of DPO training and coarse topical matching. The second objective is to determine whether the effectiveness of BASiS extends beyond a particular dialogue setting or corpus pair. We evaluate emotional support, mathematics tutoring, and medical history-taking using a different combination of small, high-quality and large general-purpose corpora for each domain. The human evaluation tests whether raters regard the corpus-specific response policies as desirable, while an LLM-based oracle evaluation on 100 dialogue contexts examines the breadth of the observed trends and differences among training conditions.

4.2 Experimental Method

We use the three corpus pairs shown in Table 1. In each setting, the small, high-quality corpus serves as the reference from which the target dialogue progression is extracted, while the large general-purpose corpus serves as the candidate pool for BASiS selection. For emotional support, we use ESConv [6], which contains dialogues between support

Table 1. Corpus pairs used for evaluation and the role of each corpus

Dialogue purpose	Small high-quality corpus	Large general-purpose corpus
Emotional support	ESConv	DailyDialog
Mathematics tutoring	MathDial	WildChat-1M
Medical history-taking	MediTOD	WildChat-1M

seekers and supporters, as the small, high-quality corpus, and DailyDialog [5], a collection of everyday multi-turn conversations, as the large general-purpose corpus. For mathematics tutoring, we use MathDial [8], which contains one-to-one mathematics dialogues between teachers and LLM-simulated learners, as the small, high-quality corpus, and WildChat-1M [20], which contains real-world interactions between users and LLMs, as the large general-purpose corpus. For medical history-taking, we use MediTOD [14], which contains clinician-authored history-taking dialogues between doctors and patients, as the small, high-quality corpus and, as in the MathDial experiment, WildChat-1M as the large general-purpose corpus.

For the BASiS condition in each setting, we first construct a Bayesian dialogue model from the small, high-quality corpus and use the Dialogue Progression Scoring Module to evaluate candidate responses in the large general-purpose corpus. We construct preference data from the resulting progression-compatibility scores and train a local LLM with LoRA-based DPO. As comparison conditions, we construct Target-only, trained only on the small target corpus, and Random, trained on data sampled from a coarsely matched topic category. For Random, broad keywords representing the main topics of the small corpus are used to filter the large corpus by category, after which the same amount of data as in BASiS is sampled from that category.

We conducted a human evaluation with nine members of our research laboratory. To limit response burden, we used 10 dialogue contexts per domain and compared Base, BASiS, and Random. For each context, the three model responses were shown as A–C without model names, and the response-to-letter mapping was varied so that a model did not consistently occupy the same position.

Participants rated every response on the questions in Table 4 using a seven-point scale from 1 (not at all applicable) to 7 (very applicable), then selected the most appropriate response or indicated that the responses were almost the same or could not be judged. All nine participants completed ESConv and MathDial. For MediTOD, we excluded two incomplete responses and analyzed the seven complete participants. For each axis, we averaged each participant’s ratings across the 10 contexts and treated participants as the analysis unit. We used the Friedman test with Kendall’s W as the effect size. For significant axes, paired permutation tests were conducted for all three pairs with Holm correction within the axis. Participant-level bootstrap 95% confidence intervals were computed using 2,000 resamples. Because each participant made repeated final choices, we report their counts and percentages descriptively without an additional test that assumes independence.

For the oracle evaluation, we prepared 100 evaluation dialogues for each setting and generated a response from each of the Base, Target-only, BASiS, and Random models. Using the same dialogues across all four conditions provides a paired comparison that controls for differences among dialogues. An LLM served as the evaluator, with model names hidden and the order of the four conditions randomized. We used GPT-5.4 at temperature 0. In addition to the preceding dialogue history and the response under evaluation, the oracle received the definition, high- and low-score conditions, and common scoring rubric for each axis in Table 3. The rubric defined 1–2 as clearly inappropriate or context-ignoring, 3–4 as weakly aligned with the axis, 5–6 as minimally adequate, 7–8 as good responses that clearly satisfy the axis, and

Table 2. DPO training and training data in each experimental condition

Condition	DPO training	Training data
Base	×	None
Target-only	✓	Small target corpus only
BASiS	✓	Progression-compatible data
Random	✓	Keyword-filtered random

9–10 as very good responses with little room for improvement; a score of 10 was reserved for a near-ideal response. The prompt instructed the evaluator not to judge responses solely by model identity or response length and to output integer scores from 1 to 10 for each axis, an overall score, and a one- or two-sentence rationale. For each oracle axis, we tested differences among the four conditions using the Friedman test. When the omnibus test was significant, we conducted paired permutation tests for all six pairwise comparisons and applied Holm correction within that axis. We used a significance level of $\alpha = .05$ and report bootstrap 95% confidence intervals for the means.

4.3 Experimental Conditions and Evaluation Measures

We conduct three comparisons to determine whether direct training on the small target corpus is sufficient, whether training with BASiS-selected data is effective, and whether the selection procedure itself contributes to performance.

Direct training on the target corpus: We compare Base with Target-only to test whether direct training on the small, high-quality target corpus is sufficient to reproduce its dialogue progression. Comparing Target-only with BASiS further tests the benefit of modeling the small corpus and using it to select data from a large general-purpose corpus.

Effect of training with BASiS-selected data: To determine whether training on data selected by BASiS improves reproduction of the target-corpus progression, we compare the Base condition, which uses the model before training, against the BASiS condition, which uses the model trained on BASiS-selected data.

Effect of data selection: To determine whether changes in the BASiS condition result from BASiS data selection rather than DPO training or coarse topical matching alone, we compare BASiS against Random. Broad keywords representing the main topics of the small corpus are first used to filter the large corpus by category. Random then samples the same amount of training data from that category and uses the same training procedure and hyperparameters as BASiS.

Table 2 summarizes the four conditions. Base is the model before DPO training. Target-only is trained by DPO using only the small, high-quality target corpus. BASiS uses data selected from the large general-purpose corpus by the progression-compatibility score, whereas Random uses data sampled from a category coarsely identified using broad keywords. All experiments use the same local LLM, Qwen3.5-72B. Base uses the pretrained model without additional training, whereas Target-only, BASiS, and Random use LoRA-based DPO. Within each experimental setting, BASiS and Random use the same amount and format of preference data. Their principal difference is therefore selection by progression compatibility using the Bayesian dialogue model versus random sampling within a coarsely matched topic category. We use a DPO learning rate of 5×10^{-6} and $\beta = 0.1$. LoRA uses rank $r = 8$, $\alpha = 16$, and dropout of 0.05, and training runs for one epoch. These settings are shared across all three corpus pairs and all three DPO-trained conditions.

Table 3 summarizes the seven oracle-evaluation axes used in each experimental setting.

Oracle evaluation axes for ESConv: For ESConv–DailyDialog, the axes capture the response policy and progression of emotional-support dialogue. In addition to the strength and tone of support characteristic of ESConv, they assess

Table 3. Oracle evaluation axes for each experiment (1–10; higher is better for every axis)

Evaluation axis	What the axis measures
Corpus pair: ESConv × DailyDialog	
ESConv Support Behavior Strength	Strength of the support behavior characteristic of ESConv
ESConv Tone Similarity	Similarity to ESConv’s empathetic, accepting, and non-assertive tone
Supporter Role Consistency	Supportive role without lecturing, self-disclosure, or digression
Non-directive Support	Respect for the seeker without assertions or imposed directions
Strategy/Stage Alignment	Fit of the response strategy to the current dialogue stage
Premature Advice Avoidance	Advice only after understanding the person’s emotions and situation
Naturalness	Natural, context-appropriate dialogue
Corpus pair: MathDial × WildChat-1M	
Equitable Tutoring	Space for the learner to think, explain, and solve independently
Learner Reasoning Diagnosis	Recognition of correctness, gaps, and confusion in learner reasoning
Mistake Location and Targeting	Accurate localization of errors or incomplete reasoning steps
Guidance Quality	Accuracy and usefulness of questions, hints, explanations, and examples
Feedback Actionability	Clarity about what the learner should think, calculate, or check next
Answer Revealing Calibration	Appropriate amount and timing of answer revelation
Teacher Move–Stage Alignment	Alignment of the teacher’s response strategy with the dialogue stage
Corpus pair: MediTOD × WildChat-1M	
Response Relevance	Relevance to the patient’s symptoms and preceding context
Overall Quality	Overall relevance, clarity, naturalness, and usefulness
Premature Assessment Avoidance	No diagnosis or treatment planning before sufficient information
Appropriate Uncertainty	Appropriate expression of uncertainty and need for confirmation
Understandability	Clear, patient-friendly questions and explanations
Unsafe Medical Advice Avoidance	Avoidance of potentially harmful medical advice
Unsupported Diagnosis Avoidance	Avoidance of unsupported disease or symptom-cause claims

maintenance of the supporter role, a non-directive response policy, alignment between response strategy and dialogue stage, the timing of advice, and naturalness as Japanese dialogue [6, 9, 19].

Oracle evaluation axes for MathDial: For MathDial–WildChat-1M, the axes evaluate not only mathematical correctness, but also whether a response encourages the learner’s own reasoning and provides pedagogical support appropriate to the learner’s current understanding [8]. They assess whether the learner has room to think, whether the system accurately diagnoses the learner’s reasoning and locates errors, and whether it selects among questions, hints, explanations, and answer revelation according to the dialogue stage.

Oracle evaluation axes for MediTOD: For MediTOD–WildChat-1M, the axes capture the appropriateness of medical history-taking, response quality, and medical safety [14]. They assess relevance to the patient’s preceding symptoms and context, overall response quality, avoidance of premature assessment, appropriate uncertainty, understandability to the patient, and avoidance of unsupported diagnosis or unsafe medical advice. Unsafe Medical Advice Avoidance and Unsupported Diagnosis Avoidance measure how well a response avoids unsafe advice or unsupported diagnosis, rather than how often it produces them. As with the other axes, higher scores therefore indicate more desirable responses.

Human-evaluation items: The human evaluation uses questions that make the response policies and dialogue progression in each domain readily interpretable to raters. For every item, 1 means “not at all applicable” and 7 means “very applicable.” Participants were not told the model names or training procedures and judged only the responses.

Table 4. Human-evaluation items for each experiment (seven-point scale; 1 = not at all applicable, 7 = very applicable)

Item	Question
Corpus pair: ESConv × DailyDialog	
Q1	Overall, this is a good response for supporting the help seeker.
Q2	The response acknowledges the help seeker’s feelings and speaks gently.
Q3	The response maintains a supportive stance that is consistent with the preceding conversation.
Q4	The response seeks to understand the person without imposing instructions or conclusions.
Q5	The timing of advice or suggestions is appropriate for this stage of the conversation.
Q6	The response fits well with what has been discussed so far.
Q7	The response is natural and easy to read as dialogue.
Final choice	Select the response that best acknowledges distress and supports continued conversation.
Corpus pair: MathDial × WildChat-1M	
Q1	The response leaves room for the learner to think, explain, and try a solution independently.
Q2	The response accurately identifies whether the learner’s reasoning is correct and where it is wrong or incomplete.
Q3	The response specifically narrows down what the learner should review or consider next.
Q4	The questions, hints, or explanations are accurate and match the learner’s current understanding.
Q5	The learner can tell what to think about, calculate, or check next.
Q6	The amount and timing of the answer or solution revealed are appropriate for the current stage of the dialogue.
Q7	The use of questions, hints, explanations, and comprehension checks is appropriate for the dialogue stage.
Final choice	Select the tutoring response that best helps the learner progress independently.
Corpus pair: MediTOD × WildChat-1M	
Q1	The response gathers necessary new information without repeating questions already answered.
Q2	The response does not rush to a diagnosis or action while information is still insufficient.
Q3	The response does not assert what remains unknown and appropriately indicates what still needs to be confirmed.
Q4	The response does not strongly prescribe medication, treatment, or care-seeking when information is insufficient.
Q5	The response does not assert a diagnosis or cause of symptoms without sufficient grounds.
Final choice	Select the clinician response that best gathers information and avoids premature conclusions.

The final-choice options were common to all domains: Response A, Response B, Response C, Almost the same, and Cannot judge.

4.4 Results

This section presents the human- and oracle-evaluation results for each experimental setting.

Experimental Setting 1: ESConv × DailyDialog

Figure 3 presents the human-evaluation results for Experimental Setting 1. BASiS had the highest mean on all seven axes and significantly exceeded both Base and Random on six, with advice timing as the exception (Holm-adjusted $p < .05$). ESConv Support Behavior Strength averaged 4.22, 5.64, and 3.97 for Base, BASiS, and Random, respectively; ESConv Tone Similarity averaged 4.59, 6.17, and 4.43. Kendall’s W ranged from .479 to .876 on significant axes. Premature Advice Avoidance also favored BASiS in the mean, but its post-hoc comparisons were nonsignificant. BASiS accounted for 74.4% of final choices.

In the oracle evaluation for Experimental Setting 1 shown in Figure 4, BASiS likewise had the highest mean on all seven axes. For Base, Target-only, BASiS, and Random, ESConv Support Behavior Strength averaged 8.26, 8.22, 8.65, and 8.21, and ESConv Tone Similarity averaged 8.21, 8.25, 8.52, and 8.13; BASiS significantly exceeded all three conditions on both axes. Premature Advice Avoidance averaged 8.57, 8.76, 9.44, and 8.92, again with significant differences from all




Fig. 3. Human-evaluation results for ESConv-DailyDialog. Bars show means of participant-level averages over 10 contexts; error bars show participant-level bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over three pairs ($*p < .05$, $**p < .01$).




Fig. 4. Oracle-evaluation results for ESConv-DailyDialog. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs ($*p < .05$, $**p < .01$, $***p < .001$).

three conditions. Naturalness was significantly higher than Random, while Target-only did not significantly exceed Base on any axis.

Experimental Setting 2: MathDial $\times$ WildChat-1M

Figure 5 presents the human-evaluation results for Experimental Setting 2. BASiS significantly exceeded both Base and Random on all seven axes (Holm-adjusted $p < .05$). Equitable Tutoring averaged 4.28, 5.42, and 3.73, and Feedback Actionability averaged 4.13, 5.43, and 3.72. Participants therefore favored BASiS both for preserving room for learner reasoning and for clarifying the next action. Kendall's W ranged from .716 to .939. BASiS accounted for 61.1% of final choices.

In the oracle evaluation for Experimental Setting 2 shown in Figure 6, BASiS exceeded all three comparison conditions on all seven axes; all 21 comparisons were significant (all $p < .001$). For Base, Target-only, BASiS, and Random, Equitable Tutoring averaged 6.77, 7.02, 7.75, and 6.00. Learner Reasoning Diagnosis averaged 7.38, 7.76, 8.70, and 6.97, and Mistake Location and Targeting averaged 7.53, 7.74, 8.85, and 7.27. The 1.85-point BASiS–Random difference on Feedback Actionability was the largest among the seven axes. Target-only had a higher mean than Base on all axes, but none of the differences was significant.

Experimental Setting 3: MediTOD $\times$ WildChat-1M

Figure 7 presents the human-evaluation results for Experimental Setting 3. BASiS had the highest mean on all five axes. It significantly exceeded both Base and Random on Premature Assessment Avoidance, Unsafe Medical Advice Avoidance, and Unsupported Diagnosis Avoidance (Holm-adjusted $p < .05$), with Kendall's W from .799 to .813. The Friedman tests were nonsignificant for Coverage Without Redundancy and Appropriate Uncertainty. BASiS was selected most often in the final choice, at 44.3%.

In the oracle evaluation for Experimental Setting 3 shown in Figure 8, BASiS likewise had the highest mean on all seven axes. For Base, Target-only, BASiS, and Random, Response Relevance averaged 4.46, 4.70, 5.09, and 4.91, and Overall Quality averaged 5.16, 5.26, 5.69, and 5.46. Appropriate Uncertainty averaged 7.84, 7.96, 8.26, and 7.92, with




Fig. 5. Human-evaluation results for MathDial-WildChat-1M. Bars and error bars show participant-level means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction ($*p < .05$, $**p < .01$).




Fig. 6. Oracle-evaluation results for MathDial-WildChat-1M. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs (all $***p < .001$).

BASiS significantly exceeding all three conditions. Unsafe Medical Advice Avoidance was significantly higher than Random, and Unsupported Diagnosis Avoidance was significantly higher than Base and Random. Target-only did not significantly exceed Base on any MediTOD axis.

Table 5 summarizes the final-choice distribution. BASiS was selected most often in every domain: 74.4% for ESConv, 61.1% for MathDial, and 44.3% for MediTOD. In MediTOD, this share was not a majority but exceeded Random by 17.2 percentage points.

4.5 Discussion

We evaluated BASiS in three distinct dialogue domains: emotional support, mathematics tutoring, and medical history-taking. BASiS had the highest mean on all 19 human-rated axes, significantly exceeded Base and Random on 16 axes each, and was the most frequent final choice in every domain. It also had the highest mean on all 21 oracle axes and significantly exceeded Base, Target-only, and Random on 18, 13, and 15 axes, respectively. We next discuss the strengths of BASiS revealed by these results and the reasons for the observed improvements.

Discussion 1: Reproducing response policies in emotional-support dialogue

Human raters preferred BASiS over Base and Random on support-behavior strength, an empathetic and accepting tone, supporter-role consistency, non-directiveness, contextual fit, and naturalness; BASiS also received 74.4% of final choices. Because both local expression and the broader stance of acknowledging feelings without imposing a solution were rated highly, BASiS appears to reproduce ESConv’s supporter behavior as an integrated response policy. One explanation is that selecting data compatible with observation labels such as `reflect_validate` and transitions toward empathetic support teaches the model to maintain a supportive role, rather than merely increasing the use of gentle wording. Premature Advice Avoidance also had the highest human mean under BASiS, while the oracle evaluation showed significant differences from all three comparison conditions. Although the human post-hoc comparisons were nonsignificant, the mean trends agree across the two evaluations. Deferring advice therefore appears as a characteristic of BASiS that could be measured more precisely with additional dialogue contexts.




Fig. 7. Human-evaluation results for MediTOD-WildChat-1M. Bars and error bars show means for the seven complete participants and participant-level bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction ($*p < .05$).




Fig. 8. Oracle-evaluation results for MediTOD-WildChat-1M. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs ($*p < .05$, $**p < .01$).

Table 5. Counts and percentages of final choices in the human evaluation

Choice	ESConv	MathDial	MediTOD
Base	14 (15.6%)	16 (17.8%)	11 (15.7%)
BASiS	67 (74.4%)	55 (61.1%)	31 (44.3%)
Random	6 (6.7%)	12 (13.3%)	19 (27.1%)
Almost the same	2 (2.2%)	3 (3.3%)	8 (11.4%)
Cannot judge	1 (1.1%)	4 (4.4%)	1 (1.4%)

Discussion 2: Reproducing pedagogical dialogue policies in mathematics tutoring

Human raters preferred BASiS over Base and Random on all seven axes, including encouraging learner reasoning, diagnosing reasoning, locating mistakes, providing useful support, calibrating answer revelation, and matching teacher moves to the dialogue stage. BASiS also received 61.1% of final choices, indicating that participants regarded the responses as desirable mathematics tutoring overall. The oracle evaluation likewise found significant differences from all three comparison conditions on every axis, making this the most consistent pattern among the three domains. In particular, gains in Feedback Actionability indicate that BASiS did not merely generate correct content, but helped learners identify what to think about, calculate, or verify next. BASiS selects examples whose transitions from the learner’s inferred state to questions, hints, or explanations match those in MathDial. This selection likely exposed the model to pedagogical scaffolding that moves a learner to the next reasoning step, explaining the broad gains from Equitable Tutoring through Teacher Move–Stage Alignment.

Discussion 3: Response quality and safety in medical history-taking

Human raters preferred BASiS over Base and Random on avoiding premature assessment, unsafe medical advice, and unsupported diagnosis. BASiS was also the most frequent final choice at 44.3%, 17.2 percentage points above Random. Its clearest strength in medical history-taking therefore appears to be a cautious response policy that withholds

conclusions until sufficient information is available. By evaluating candidates through state transitions, BASiS can select responses that gather confirmation before moving to assessment, rather than responses that merely contain medically related words. Coverage Without Redundancy and Appropriate Uncertainty did not show significant human condition differences, although BASiS had the highest mean on both. The oracle evaluation also placed BASiS highest on all seven axes and found a significant advantage over all three conditions on Appropriate Uncertainty. Together, these results suggest that the benefits are clearest for safety-critical forms of restraint, while the efficiency of information gathering and the expression of uncertainty may vary more with the dialogue context. More varied symptoms and dialogue stages would allow these aspects to be examined in greater detail.

Discussion 4: Selecting dialogue progression rather than response content alone

Across all three domains, the strongest results concerned not only topic or tone, but also response timing and dialogue progression. Human raters valued acknowledging feelings in emotional support, scaffolding the learner’s next step in mathematics tutoring, and avoiding premature conclusions in medical history-taking. Each policy requires the model to consider both the preceding history and the current dialogue stage. This pattern accords with Equation (3). BASiS does not score candidates solely by generic single-turn quality; it uses posterior probabilities based on state transitions and observation labels. Consequently, knowledge about *when* to offer a given kind of response is reflected in selection and appears in generated responses that humans judge desirable across the dialogue as a whole.

Discussion 5: Contributions of the target corpus and data selection

In the oracle evaluation, Target-only did not significantly exceed Base on any of the 21 axes, whereas BASiS exceeded Target-only on 13. A small target corpus can therefore be used more effectively by extracting its progression patterns and treating them as selection criteria for a large corpus, rather than using it only as direct training data. BASiS also had a higher oracle mean than Random on all 21 axes and exceeded Random on 16 of the 19 human axes. Both conditions used the same amount of training data and the same training configuration, and Random was already restricted to a coarsely matched topic category. The consistent differences therefore indicate that transferring corpus-specific dialogue progression benefits from selecting examples aligned with conversational states, response strategies, and state transitions in addition to broad topical similarity.

Discussion 6: Applicability across three dialogue domains

We applied the same BASiS pipeline to emotional support, mathematics tutoring, and medical history-taking by changing only the reference and selection corpora. BASiS had the highest mean on every human and oracle axis and was the most frequent final choice in every domain. It can therefore extract different corpus-specific response policies rather than being limited to a single tone. The fact that the strongest gains differed by domain further illustrates the advantage of inducing states and strategies from each corpus instead of imposing one fixed guideline. Thus, although these findings provide clear initial evidence for the effectiveness of BASiS, generalization to broader user populations requires validation with external raters and more dialogue contexts. Future work should broaden the scope of evaluation and refine the design of corpus-specific states, strategies, and desirable state sets, thereby extending the strengths of BASiS to a wider range of dialogue applications.

5 Conclusion

We proposed BASiS (Bayesian Style-based Selection), a method that transfers dialogue progression from a small, high-quality corpus to a local large language model (LLM) through Bayesian data selection. Through constrained ontology induction, BASiS derives conversational states, response strategies, and state transitions from the small corpus

and represents them as a Bayesian dialogue model. It uses this model to evaluate candidate responses in a large general-purpose corpus and trains a local LLM on the resulting preference data. This mechanism uses the response policies and multi-turn progression captured in a small corpus as criteria for selecting training data, without requiring detailed dialogue guidelines to be written manually for each setting. By representing conversational states as a probability distribution and updating that distribution with observations from candidate responses, BASiS selects data according to the dialogue history and the state transition induced by each response.

In comparative experiments across emotional support, mathematics tutoring, and medical history-taking, BASiS achieved the highest mean on all 19 human-rated axes and was the most frequent final choice in every domain. It also achieved the highest mean on all 21 oracle axes. Together, the evaluations suggest that BASiS can incorporate empathetic support, pedagogical guidance suited to learner understanding, and response policies that avoid unwarranted certainty when information is insufficient. In the oracle evaluation, Target-only did not significantly outperform Base on any axis, and BASiS broadly outperformed Random. These findings demonstrate the value of using the small corpus as a source of interactional knowledge and selecting training data beyond coarse topical similarity.

Future work will evaluate other models, investigate how to design conversational states, response strategies, and desirable state sets for different dialogue domains, and conduct human evaluations with external raters and more dialogue contexts. Through these efforts, we aim to extend BASiS across applications including counseling, education, and medicine.

References

- [1] Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, and Hongxia Jin. 2024. AlpaGasus: Training a Better Alpaca with Fewer Data. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=FdVXgSJhvz>
- [2] Ananya Ganesh, Martha Palmer, and Katharina Kann. 2021. What Would a Teacher Do? Predicting Future Talk Moves. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*. Association for Computational Linguistics, Online, 4739–4751. doi:10.18653/v1/2021.findings-acl.418
- [3] Prakhar Gupta, Yang Liu, Di Jin, Behnam Hedayatnia, Spandana Gella, Sijia Liu, Patrick Lange, Julia Hirschberg, and Dilek Hakkani-Tur. 2023. DialGuide: Aligning Dialogue Model Behavior with Developer Guidelines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 14031–14047. doi:10.18653/v1/2023.findings-emnlp.935
- [4] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. The Tenth International Conference on Learning Representations (ICLR 2022). <https://openreview.net/forum?id=nZcVKeeFYf9>
- [5] Yanran Li, Hui Su, Xiaoyu Shen, Wenjie Li, Ziqiang Cao, and Shuzi Niu. 2017. DailyDialog: A Manually Labelled Multi-turn Dialogue Dataset. In *Proceedings of the Eighth International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Asian Federation of Natural Language Processing, Taipei, Taiwan, 986–995. <https://aclanthology.org/I17-1099/>
- [6] Siyang Liu, Chujie Zheng, Orianna Demasi, Sahand Sabour, Yu Li, Zhou Yu, Yong Jiang, and Minlie Huang. 2021. Towards Emotional Support Dialog Systems. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Association for Computational Linguistics, Online, 3469–3483. doi:10.18653/v1/2021.acl-long.269
- [7] Wei Liu, Weihao Zeng, Keqing He, Yong Jiang, and Junxian He. 2024. What Makes Good Data for Alignment? A Comprehensive Study of Automatic Data Selection in Instruction Tuning. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=BTKAeLqLMw>
- [8] Jakub Macina, Nico Daheim, Sankalan Pal Chowdhury, Tanmay Sinha, Manu Kapur, Iryna Gurevych, and Mrinmaya Sachan. 2023. MathDial: A Dialogue Tutoring Dataset with Rich Pedagogical Properties Grounded in Math Reasoning Problems. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 5602–5621. doi:10.18653/v1/2023.findings-emnlp.372
- [9] Remi Mir, Bjarke Felbo, Nick Obradovich, and Iyad Rahwan. 2019. Evaluating Style Transfer for Text. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, 495–504. doi:10.18653/v1/N19-1049
- [10] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems*, Vol. 35. Curran Associates, Inc., Red Hook, NY, USA, 27730–27744.

- [11] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 53728–53741. doi:10.52202/075280-2338
- [12] Hannah Rashkin, Eric Michael Smith, Margaret Li, and Y-Lan Boureau. 2019. Towards Empathetic Open-domain Conversation Models: A New Benchmark and Dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, Florence, Italy, 5370–5381. doi:10.18653/v1/P19-1534
- [13] Eric Rudolph, Philipp Steigerwald, and Jens Albrecht. 2026. Transition-Matrix Regularization for Next Dialogue Act Prediction in Counselling Conversations. In *Findings of the Association for Computational Linguistics: ACL 2026*. Association for Computational Linguistics, San Diego, California, United States, 25442–25457. doi:10.18653/v1/2026.findings-acl.1271
- [14] Vishal Vivek Saley, Goonjan Saha, Rocktim Jyoti Das, Dinesh Raghu, and Mausam. 2024. MediTOD: An English Dialogue Dataset for Medical History Taking with Comprehensive Annotations. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Miami, Florida, USA, 16843–16877. doi:10.18653/v1/2024.emnlp-main.936
- [15] Chenwei Wan, Matthieu Labeau, and Chloé Clavel. 2025. EmoDynamix: Emotional Support Dialogue Strategy Prediction by Modelling MiXed Emotions and Discourse Dynamics. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics, Albuquerque, New Mexico, 1678–1695. doi:10.18653/v1/2025.naacl-long.81
- [16] Xiangyu Wen, Jianyuan Zhong, Zhijian Xu, and Qiang Xu. 2025. Guideline Compliance in Task-Oriented Dialogue: The Chained Prior Approach. In *Findings of the Association for Computational Linguistics: NAACL 2025*. Association for Computational Linguistics, Albuquerque, New Mexico, 6765–6791. doi:10.18653/v1/2025.findings-naacl.377
- [17] Mengzhou Xia, Sadhika Malladi, Suchin Gururangan, Sanjeev Arora, and Danqi Chen. 2024. LESS: Selecting Influential Data for Targeted Instruction Tuning. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, Vienna, Austria, 54104–54132. <https://proceedings.mlr.press/v235/xia24c.html>
- [18] Bolin Zhang, Jiahao Wang, Qianlong Du, Jiajun Zhang, Zhiying Tu, and Dianhui Chu. 2025. A Survey on Data Selection for LLM Instruction Tuning. *Journal of Artificial Intelligence Research* 83, Article 32 (2025), 14 pages. doi:10.1613/jair.1.17625
- [19] Saizheng Zhang, Emily Dinan, Jack Urbanek, Arthur Szlam, Douwe Kiela, and Jason Weston. 2018. Personalizing Dialogue Agents: I have a dog, do you have pets too?. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Melbourne, Australia, 2204–2213. doi:10.18653/v1/P18-1205
- [20] Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. 2024. WildChat: 1M ChatGPT Interaction Logs in the Wild. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=Bl8u7ZRIbM>
- [21] Chunting Zhou, Pengfei Liu, Puxin Xu, Srinu Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, Lili Yu, Susan Zhang, Gargi Ghosh, Mike Lewis, Luke Zettlemoyer, and Omer Levy. 2023. LIMA: Less Is More for Alignment. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 55006–55021. doi:10.52202/075280-2400