

BASiS: Transferring Dialogue Progression Strategies to Large Language Models through Bayesian Dialogue Modeling

ANONYMOUS AUTHOR(S)

We propose BASiS (Bayesian Style-based Selection), a framework for scaling the interactional knowledge encoded in a small, high-quality dialogue corpus to large-scale LLM training. We focus on settings where dialogue must progress according to the user’s state and the stage of the interaction as exemplified in emotional support, education, and healthcare. In such settings, dialogue quality depends not only on utterance content but also on how the dialogue develops across turns, which involves procedural interactional knowledge of how to adapt response strategies as the interaction evolves. Although such knowledge is implicitly encoded in small, high-quality corpora, these corpora alone provide limited data for LLM training. To overcome this challenge, BASiS models the progression patterns in the high-quality corpus and uses the resulting model to guide training-data selection from a large general-purpose corpus. It extracts dialogue states, response strategies, and state transitions from the small corpus and represents the probabilistic relationships as a Bayesian dialogue model. The model scores candidate responses in a large corpus according to how well they support the dialogue progression exhibited in the high-quality corpus, and the scores are used to construct preference pairs for LLM fine-tuning with Direct Preference Optimization (DPO). We evaluated BASiS in emotional support, mathematics tutoring, and medical history-taking against the base model and a model trained on randomly selected in-category data. In the human evaluation, BASiS achieved the highest mean score on all 19 evaluation dimensions and was selected most often in all three domains. In the LLM-based oracle evaluation, BASiS achieved the highest mean score on all 21 evaluation dimensions, with broad statistically significant improvements in emotional support and mathematics tutoring.

CCS Concepts: • **Computing methodologies** → **Natural language generation**; **Discourse, dialogue and pragmatics**; Learning from implicit feedback; • **Human-centered computing** → *Natural language interfaces*.

Additional Key Words and Phrases: dialogue progression, data selection, Bayesian dialogue modeling, response strategy, preference optimization, emotional support dialogue, large language models

ACM Reference Format:

Anonymous Author(s). 2026. BASiS: Transferring Dialogue Progression Strategies to Large Language Models through Bayesian Dialogue Modeling. In *Proceedings of 32nd Annual ACM Conference on Intelligent User Interfaces (IUI '27)*. ACM, New York, NY, USA, 17 pages.

1 Introduction

Advances in instruction tuning and reinforcement learning from human feedback (RLHF) have enabled recent large language models (LLMs) to follow instructions, generate accurate responses, and sustain natural and coherent multi-turn conversations across a wide range of topics [10]. In settings such as emotional support, mathematics tutoring, and medical history-taking, however, a system must do more than provide correct information: it must strategically progress the dialogue flow according to the interlocutor’s state and the current stage of the interaction [6, 12]. In mathematics tutoring, for example, the system should not simply provide the correct answer; it must first understand

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

the learner’s current reasoning and then choose a question, hint, or explanation that helps the learner take the next step independently. Producing such responses requires the system to control dialogue progression across turns.

LLM-based response generation has approached this requirement in two main ways. Guideline-based methods describe desirable response strategies in natural-language dialogue guidelines and guide generation to follow them [3, 16]. Data-driven methods instead adapt model behavior through instruction tuning or by selecting high-quality training examples [18, 21]. However, the former requires dialogue guidelines to be designed manually for each setting, making it difficult to enumerate implicit features such as response timing and multi-turn progression. The latter depends on sufficient training data, whereas detailed interactional knowledge is often encoded in small, carefully constructed dialogue corpora that alone provide too few examples for LLM training.

A separate line of research has modeled dialogue progression through dialogue actions, response strategies, and their transitions. Ganesh et al. predicted the next teacher talk move in classroom dialogue [2]; Wan et al. modeled relationships between user emotion states and response strategies as a heterogeneous graph [15]; and Rudolph et al. regularized next-dialogue-act prediction using a transition matrix derived from counseling conversations [13]. These studies capture important structure beyond surface utterances, but primarily evaluate action or strategy prediction. Their integration with LLMs for natural-language response generation, and evaluation of whether the modeled progression appears in generated responses, remain limited.

To bridge these two lines of work, we propose BASiS (Bayesian Style-based Selection), a method for transferring dialogue progression from a small, high-quality corpus to an LLM. BASiS first induces a progression model comprising conversational states, response strategies, state transitions, and their probabilistic relationships. It then scores candidate responses in a large general-purpose corpus according to their compatibility with that progression and constructs preference pairs from the scores. Finally, it adapts the LLM using Direct Preference Optimization (DPO) [11] with Low-Rank Adaptation (LoRA) [4]. This combination uses large-corpus selection to scale the interactional knowledge encoded in the small corpus, while reducing the need to manually author detailed dialogue guidelines. By maintaining a probability distribution over conversational states, BASiS can account for dialogue history and multiple valid response strategies. We evaluate the same pipeline in emotional support, mathematics tutoring, and medical history-taking by changing the reference corpus and candidate pool.

The remainder of this paper is organized as follows. Section 2 reviews related work on guideline-based response control, dialogue-history-based prediction of dialogue actions and response strategies, and training-data selection for LLMs. Section 3 describes the four steps of the BASiS pipeline. Section 4 presents the experiments used to evaluate BASiS, their results, and our discussion. Finally, Section 5 concludes the paper.

2 Related Work

Guideline-based approaches make desired dialogue behavior explicit and use it to control response generation. DialGuide represents each guideline as a natural-language condition specifying when it applies and an action specifying what the response should contain [3]. It retrieves guidelines relevant to the current dialogue context and uses them to generate or verify a response. GuidedTOD instead represents actions specified in operational guidelines as states in a Chained Prior, a Markov chain whose transition probabilities are calibrated from curated service data [16]. During inference, the Chained Prior is updated as the dialogue progresses and is used to rerank candidate outputs. These studies show that explicit dialogue guidelines can guide generation and selection according to the current stage of a dialogue. However, developers must write the guidelines and redesign them when the desired dialogue progression changes.

Dialogue progression can also be modeled from corpora. Ganesh et al. predict the teacher’s next talk move from the preceding classroom dialogue and talk moves [2]. EmoDynamix represents discourse-level relationships between fine-grained user emotion states and system response strategies as a heterogeneous graph [15]. Rudolph et al. estimate dialogue-act transition patterns from counseling conversations and use them to regularize next-dialogue-act prediction toward the observed dialogue flow [13]. Together, these studies establish that dialogue progression can be represented as corpus-derived states, actions, strategies, and transitions rather than only as manually written rules.

Data-centric approaches select training examples from a large candidate pool to adapt LLM behavior efficiently. LIMA emphasizes careful curation for data quality and diversity [21]. AlpaGasus filters instruction data using LLM-based quality scores, whereas DEITA combines instruction complexity, response quality, and diversity-aware selection [1, 7]. LESS estimates the influence of individual examples on downstream performance and selects those expected to be most useful [17]. These studies show that a criterion representing the desired behavior or capability can guide the construction of effective adaptation data.

Corpus-based dialogue modeling and data-centric adaptation together suggest a broader opportunity: using dialogue progression extracted from a corpus as a criterion for selecting LLM training data. In settings such as mathematics tutoring, the suitability of a response depends not only on its standalone quality or topical relevance, but also on the learner’s current state, the preceding interaction, and the stage toward which the response moves the dialogue. The corpus-based models above use dialogue structure primarily to predict the next action or strategy rather than to supervise natural-language response generation. Conversely, the data-selection methods above assess examples by quality, complexity, diversity, or downstream influence without representing how each response advances a multi-turn interaction as learned from a high-quality dialogue corpus. BASiS addresses this gap by converting the implicit dialogue progression in such a corpus into a probabilistic selection criterion and then into preference data for LLM adaptation. In this way, BASiS enables corpus-derived interactional knowledge to shape LLM-generated responses at scale, without requiring developers to specify the complete dialogue progression through manually authored guidelines.

3 BASiS

3.1 Concept

We propose BASiS, which adapts an LLM to the response strategies and dialogue progression encoded in a small, high-quality corpus through Bayesian data selection. BASiS converts these interactional patterns into a probabilistic dialogue model. Instead of training directly on the small corpus, it uses the model as a criterion for selecting data from a large general-purpose corpus. BASiS evaluates each dialogue instance in the large corpus according to its compatibility with the target corpus’s dialogue progression. It thereby incorporates the interactional knowledge in the small corpus into training data without requiring detailed dialogue guidelines to be written manually for each setting. The selected data are ultimately used for DPO training, which adapts the LLM to the response strategies and progression patterns of the small corpus.

3.2 BASiS Pipeline

Figure 1 presents the four steps of the BASiS pipeline, from modeling dialogue progression to selecting training data and adapting an LLM. The upper pathway A, *Dialogue Progression Modeling*, comprises Steps 1 and 2. Step 1 takes the dialogue text and existing annotations in a small, high-quality corpus as input and extracts corpus-specific conversational states, response strategies, state transitions, and probability estimates. Step 2 takes these outputs, normalizes the probability




Fig. 1. BASiS pipeline. The upper pathway extracts dialogue features and probability estimates from a small, high-quality corpus and formalizes them as a Bayesian dialogue model. The lower pathway applies the fixed model to candidate responses from a large general-purpose corpus and uses the resulting preference data to train a LoRA adapter with DPO. The numbered circles correspond to the four steps described in the text.

estimates, and organizes them as a fixed Bayesian dialogue model. The lower pathway B, *Data Selection & Adaptation*, comprises Steps 3 and 4. Step 3 takes the fixed Bayesian model and dialogue data from a large general-purpose candidate pool as input and assigns each candidate response a progression-compatibility score between 0 and 1. Step 4 takes the candidate responses and their scores, constructs *chosen-rejected* preference pairs, and outputs a LoRA adapter trained with DPO. Applying this adapter to the base LLM yields a model adapted to the progression patterns in the reference corpus. Thus, the small corpus serves as the progression reference, whereas the large corpus provides the candidate training data. The same pipeline can be applied in other dialogue settings by changing these two corpora.

Step 1: Dialogue Feature Analysis

The input to Step 1 is the dialogue text and existing annotations for 80 conversations selected from a small, high-quality corpus. All 80 conversations are supplied together in a single prompt rather than divided across calls. The output is a representation of the corpus comprising conversational states, response strategies, state transitions, and probability estimates. For each dialogue turn, a *conversational state* describes the interlocutor’s situation or the current stage of the interaction, a *response strategy* describes the observable communicative function of the response, and a *state transition* describes the change from the state before the response to the state after it.

The prompt instructs the LLM to consolidate semantically similar labels across the selected dialogues, define 4–7 conversational states, keep states conceptually distinct from response strategies, assign these labels to the turns, and identify which states represent desirable progression. Existing annotations provide evidence about the communicative role of each turn, but their labels are not mechanically copied as state names. Rather than calculating probabilities from label frequencies, the LLM outputs numerical estimates of the initial-state probabilities, state-transition probabilities, and observation likelihoods based on the corpus as a whole. We then check for missing labels, values outside $[0, 1]$, and rows of probabilities that do not sum to one; Step 2 normalizes the latter where necessary. For example, ESConv, the emotional-support dialogue corpus used in our evaluation [6], yielded *opening_rapport*, a state in which greetings and brief check-ins create a safe atmosphere, and *empathic_support*, a state in which the help seeker’s emotions are acknowledged to alleviate isolation and distress. Its response strategies included *explore_question*, which asks




Fig. 2. Schematic example of candidate-response evaluation in Step 3. Panels (a)–(e) show the five internal operations: prior-state specification, prediction, response labeling, Bayesian update, and progression-compatibility scoring. The blue and orange paths show where the predicted distribution and observation label enter the update equation. Here, the probabilities of the target states Processing and Solutions are summed to obtain Score = 0.88.

contextually relevant questions to understand the seeker’s situation, and `reflect_validate`, which paraphrases and validates the seeker’s words and emotions.

Step 2: Bayesian Dialogue Modeling

Step 2 takes the labels and probability estimates produced in Step 1 as input and outputs a normalized, fixed Bayesian dialogue model representing the progression in the reference corpus. Let $S = \{s_1, \dots, s_K\}$ denote the conversational states and $O = \{o_1, \dots, o_L\}$ the observation labels corresponding to response strategies. The model consists of the initial-state probability $P(s_1)$, state-transition probability $P(s_t | s_{t-1})$, and observation likelihood $P(o_t | s_t)$. The transition probabilities represent the relative tendency, estimated from the selected reference dialogues, to move from s_{t-1} to s_t ; the observation likelihood represents the probability of observing strategy o_t in state s_t . Step 2 normalizes the estimates from Step 1 into valid probability distributions and fixes them as the initial-state distribution, transition matrix, and observation-likelihood matrix. It performs no additional corpus analysis or model fitting. The resulting model represents probabilistic relationships between states and strategies rather than memorizing utterances, so it can be applied to candidate responses outside the small corpus.

Step 3: Dialogue Progression Scoring

Step 3 takes the fixed Bayesian dialogue model from Step 2 together with dialogue histories and candidate responses from the large general-purpose corpus as input. Without rebuilding or fitting the model, it applies Bayesian inference to each candidate and outputs a progression-compatibility score. Figure 2 expands this step into five internal operations: (a) specifying the prior state distribution, (b) predicting the next-state distribution, (c) labeling the candidate response, (d) updating the state distribution, and (e) computing the score.

To construct the state distribution for each candidate, we group the large corpus by conversation ID and process responses from the beginning of each conversation in the order indicated by the turn index. The first evaluated response starts from the initial-state distribution produced in Step 2. Each preceding assistant response is assigned one of the response-strategy labels and used to update the Bayesian state distribution; the resulting posterior is carried forward as the prior for the next response. The LLM receives the complete dialogue history from the beginning of the conversation when assigning a response label. When a *chosen* response and its generated *rejected* alternatives are compared for the same context, all candidates use the same state distribution immediately before that response.

Before a candidate response is presented, the dialogue state is maintained as a probability distribution over the state set S rather than fixed to a single state (panel (a) in Figure 2). In the example, the state distribution through turn $t - 1$, $b_{t-1}(s)$, assigns 0.30 to Emotion disclosure, 0.45 to Emotion processing, and 0.25 to Considering solutions. The predicted state distribution at turn t is then computed using the fixed transition probabilities (panel (b)):

$$\hat{b}_t(s_t) = \sum_{s_{t-1} \in S} P(s_t | s_{t-1}) b_{t-1}(s_{t-1}). \quad (1)$$

The state-transition diagram in panel (b) of Figure 2 schematically shows transitions among Disclosure, Processing, and Solutions, together with representative response strategies observed around those transitions. Equation (1) maps the previous distribution b_{t-1} to the predicted distribution $\hat{b}_t$ using $P(s_t | s_{t-1})$, which summarizes the transition patterns in the reference corpus.

Next, an LLM maps candidate response r_t to exactly one observation label o_t from the fixed set of response-strategy labels (panel (c)). The input comprises the dialogue history, candidate response, and the labels with their descriptions; the LLM is not allowed to create a new label or state name. Only the selected observation label is used in the subsequent Bayesian computation. In the example, “That sounds really difficult.” is labeled *Empathy / validation*. The LLM does not assign a conversational state directly. Instead, the fixed observation likelihood for o_t is used to update the predicted distribution by Bayes’ rule:

$$b_t(s_t) = \frac{P(o_t | s_t) \hat{b}_t(s_t)}{\sum_{s' \in S} P(o_t | s') \hat{b}_t(s')}. \quad (2)$$

Here, $b_t(s_t)$ is the posterior probability $P(s_t | o_{1:t})$ conditioned on the observation history $o_{1:t}$ after incorporating the observations into the predicted distribution $\hat{b}_t(s_t)$. In panel (d) of Figure 2, the orange path shows that observation o_t determines the likelihood term $P(o_t | s_t)$, whereas the blue path shows where prediction $\hat{b}_t(s_t)$ enters the update. The denominator normalizes the posterior over all states. In this example, the updated distribution assigns 0.12 to Disclosure, 0.66 to Processing, and 0.22 to Solutions.

Finally, let $S_{\text{target}} \subseteq S$ be the set of states defined as desirable for the target dialogue progression. The progression-compatibility score of candidate response r_t is the total probability assigned to S_{target} in the updated state distribution:

$$\text{Score}(r_t) = \sum_{s \in S_{\text{target}}} b_t(s). \quad (3)$$

The progression-compatibility score ranges from 0 to 1. A larger value indicates that the state distribution after observing the candidate response is more strongly aligned with states considered desirable in the target corpus. In panel (e) of Figure 2, $S_{\text{target}} = \{\text{Processing, Solutions}\}$, so the score is $0.66 + 0.22 = 0.88$. More generally, S_{target} may include states such as processing emotions or considering appropriate solutions in emotional support, advancing the learner’s own reasoning in mathematics tutoring, or progressively organizing necessary symptom information in medical history-taking.

This scoring procedure maintains a probability distribution over states instead of fixing a single response strategy as the correct next action. Even when each candidate response receives one observation label, responses associated with different strategies can each obtain a high progression-compatibility score if those strategies have high observation likelihoods in the desirable states. Thus, when several response strategies are valid for the same context, Step 3 can evaluate them without restricting the choice to a single strategy.

Step 4: Preference Data Selection and DPO Adaptation

Step 4 takes the candidate responses and progression-compatibility scores from Step 3 as input and outputs a LoRA adapter trained with DPO. Each original context contributes one actual assistant response rather than a pre-existing set of alternatives. We first score these responses and retain high-scoring ones as *chosen* candidates. For each retained context, an LLM generates eight natural responses that are weaker in their response strategy. We do not use responses taken from other dialogues as negative examples. All eight alternatives are scored from the same state distribution immediately before the response, and the alternative with the largest score difference from the chosen response is selected as *rejected*.

We retain a BASiS preference pair only when the chosen score is at least 0.70, the rejected score is at most 0.55, and their difference is at least 0.20. The resulting pairs are used for Direct Preference Optimization (DPO) [11], which optimizes the model to prefer chosen responses over rejected ones. Rather than updating all model parameters, we apply Low-Rank Adaptation (LoRA) [4] and update only the added low-rank parameters. The resulting adapter can be attached to the base LLM, efficiently incorporating the dialogue progression represented in the selected data while retaining the model’s general language-generation capabilities.

4 Evaluation

4.1 Evaluation Objectives

Our evaluation has two objectives. The first is to determine whether a model trained on data selected by BASiS can reproduce the dialogue progression and response strategies captured in a small, high-quality corpus. We compare the model trained on BASiS-selected data against both the untrained base model and a model trained on data randomly sampled from a category coarsely filtered using keywords that describe the target corpus. These comparisons test whether BASiS training strengthens the dialogue behavior represented in the target corpus and whether the observed change is associated with BASiS data selection beyond the effects of DPO training and coarse topical matching. The second objective is to determine whether the effectiveness of BASiS extends beyond a particular dialogue setting or corpus pair. We evaluate emotional support, mathematics tutoring, and medical history-taking using a different combination of small, high-quality and large general-purpose corpora for each setting. The human evaluation tests whether raters regard the corpus-specific response strategies as desirable, while an LLM-based oracle evaluation on 100 dialogue contexts examines the breadth of the observed trends and differences among training conditions.

4.2 Experimental Method

We use the three corpus pairs shown in Table 1. In each setting, the small, high-quality corpus serves as the reference from which dialogue progression is extracted, while the large general-purpose corpus serves as the candidate pool for BASiS selection. For emotional support, we use ESConv [6], which contains dialogues between support seekers and supporters, as the small, high-quality corpus, and DailyDialog [5], a collection of everyday multi-turn conversations, as the large general-purpose corpus. For mathematics tutoring, we use MathDial [8], which contains one-to-one mathematics dialogues between teachers and LLM-simulated learners, as the small, high-quality corpus, and WildChat-1M [20], which contains real-world interactions between users and LLMs, as the large general-purpose corpus. For medical history-taking, we use MediTOD [14], which contains clinician-authored history-taking dialogues between doctors and patients, as the small, high-quality corpus and, as in the MathDial experiment, WildChat-1M as the large

Table 1. Corpus pairs used for evaluation and the role of each corpus

Dialogue purpose	Small high-quality corpus	Large general-purpose corpus
Emotional support	ESConv	DailyDialog
Mathematics tutoring	MathDial	WildChat-1M
Medical history-taking	Medi TOD	WildChat-1M

general-purpose corpus. For Step 1, we select 80 dialogues from each reference corpus. The selection covers the strategy, emotion, and problem annotations in ESConv; teacher moves in MathDial; and dialogue lengths and slots in Medi TOD.

For the BASiS condition in each setting, we construct the preference pairs using the four steps in Section 3. The pair counts are 2,500 for ESConv, 2,500 for MathDial, and 2,324 for Medi TOD. The comparison conditions are Target-only, trained using preference pairs derived from the small target corpus, and Random, trained using pairs whose chosen responses are sampled without using the BASiS score.

For Random, we first restrict the large-corpus candidate pool using domain-related terms—for example, sad, anxious, and relationship problem for ESConv; learn, tutor, and math for MathDial; and symptom, pain, and doctor for Medi TOD. We then sample actual assistant responses from the filtered pool with seed 42 and use them as chosen responses. For each context, we generate eight generally low-quality alternatives, score all of them with BASiS, and use the lowest-scoring alternative, which maximizes the difference from the chosen response, as rejected. BASiS scores are thus used only to choose the rejected response, not to select Random’s chosen responses. Random uses the same pair counts as BASiS.

For Target-only, actual assistant responses from the training split of each small corpus serve as chosen responses. We sample them while balancing ESConv strategies, MathDial teacher moves, or Medi TOD response slots. For each context, we generate and rescore eight weaker alternatives in the same way and select the alternative with the largest score difference as rejected. Target-only contains 500 preference pairs per domain.

The LLM operations within BASiS use the Responses API. We use gpt-5.4-pro for the corpus analysis and probability estimation in Step 1 and gpt-5.4 for response-strategy labeling in Step 3. Both use temperature 1.0 and reasoning effort high. The oracle evaluation uses the separate configuration described below.

We conducted a human evaluation with nine members of our research laboratory. To limit response burden, we used 10 dialogue contexts per domain and compared Base, BASiS, and Random. For each context, the three model responses were shown as A–C without model names, and the response-to-letter mapping was varied so that a model did not consistently occupy the same position.

Participants rated every response on the questions in Table 4 using a seven-point scale from 1 (not at all applicable) to 7 (very applicable), then selected the most appropriate response or indicated that the responses were almost the same or could not be judged. All nine participants completed ESConv and MathDial. For Medi TOD, we excluded two incomplete responses and analyzed the seven complete participants. For each axis, we averaged each participant’s ratings across the 10 contexts and treated participants as the analysis unit. We used the Friedman test with Kendall’s W as the effect size. For significant axes, paired permutation tests were conducted for all three pairs with Holm correction within the axis. Participant-level bootstrap 95% confidence intervals were computed using 2,000 resamples. Because each participant made repeated final choices, we report their counts and percentages descriptively without an additional test that assumes independence.

Table 2. DPO training and training data in each experimental condition

Condition	DPO training	Training data
Base	×	None
Target-only	✓	Small target corpus only
BASiS	✓	BASiS-selected data
Random	✓	Random sample after filtering

For the oracle evaluation, we prepared 100 evaluation dialogues for each setting and generated a response from each of the Base, Target-only, BASiS, and Random models. Using the same dialogues across all four conditions provides a paired comparison that controls for differences among dialogues. An LLM served as the evaluator, with model names hidden and the order of the four conditions randomized. We used GPT-5.4 at temperature 0. In addition to the preceding dialogue history and the response under evaluation, the oracle received the definition, high- and low-score conditions, and common scoring rubric for each axis in Table 3. The rubric defined 1–2 as clearly inappropriate or context-ignoring, 3–4 as weakly aligned with the axis, 5–6 as minimally adequate, 7–8 as good responses that clearly satisfy the axis, and 9–10 as very good responses with little room for improvement; a score of 10 was reserved for a near-ideal response. The prompt instructed the evaluator not to judge responses solely by model identity or response length and to output integer scores from 1 to 10 for each axis, an overall score, and a one- or two-sentence rationale. For each oracle axis, we tested differences among the four conditions using the Friedman test. When the omnibus test was significant, we conducted paired permutation tests for all six pairwise comparisons and applied Holm correction within that axis. We used a significance level of $\alpha = .05$ and report bootstrap 95% confidence intervals for the means.

4.3 Experimental Conditions and Evaluation Measures

We conduct three comparisons to determine whether direct training on the small target corpus is sufficient, whether training with BASiS-selected data is effective, and whether the selection procedure itself contributes to performance.

Direct training on the target corpus: We compare Base with Target-only to test whether direct training on the small, high-quality target corpus is sufficient to reproduce its dialogue progression. Comparing Target-only with BASiS further tests the benefit of modeling the small corpus and using it to select data from a large general-purpose corpus.

Effect of training with BASiS-selected data: To determine whether training on data selected by BASiS improves reproduction of the target corpus’s dialogue progression, we compare the Base condition, which uses the model before training, against the BASiS condition, which uses the model trained on BASiS-selected data.

Effect of data selection: To determine whether changes in the BASiS condition result from BASiS data selection rather than DPO training or coarse topical matching alone, we compare BASiS against Random. Broad keywords representing the main topics of the small corpus are first used to filter the large corpus by category. Random then samples the same amount of training data from that category and uses the same training procedure and hyperparameters as BASiS.

Table 2 summarizes the four conditions. Base is the model before DPO training. Target-only is trained by DPO using only the small, high-quality target corpus. BASiS uses data selected from the large general-purpose corpus by the compatibility score, whereas Random uses data sampled from a category coarsely identified using broad keywords. All experiments use the same LLM, Qwen3.5-27B. Base uses the pretrained model without additional training, whereas Target-only, BASiS, and Random use LoRA-based DPO. Within each experimental setting, BASiS and Random use the

Table 3. Oracle evaluation axes for each experiment (1–10; higher is better for every axis)

Evaluation axis	What the axis measures
Corpus pair: ESConv × DailyDialog	
ESConv Support Behavior Strength	Strength of support behavior characteristic of ESConv
ESConv Tone Similarity	Similarity to ESConv’s empathetic, accepting, and non-assertive tone
Supporter Role Consistency	Supportive role without lecturing, self-disclosure, or digression
Non-directive Support	Respect for the seeker without assertions or imposed directions
Strategy/Stage Alignment	Fit of the response strategy to the current dialogue stage
Premature Advice Avoidance	Advice only after understanding the person’s emotions and situation
Naturalness	Natural, context-appropriate dialogue
Corpus pair: MathDial × WildChat-1M	
Equitable Tutoring	Space for the learner to think, explain, and solve independently
Learner Reasoning Diagnosis	Recognition of correctness, gaps, and confusion in learner reasoning
Mistake Location and Targeting	Accurate localization of errors or incomplete reasoning steps
Guidance Quality	Accuracy and usefulness of questions, hints, explanations, and examples
Feedback Actionability	Clarity about what the learner should think, calculate, or check next
Answer Revealing Calibration	Appropriate amount and timing of answer revelation
Teacher Move–Stage Alignment	Alignment of the teacher’s response strategy with the dialogue stage
Corpus pair: MediTOD × WildChat-1M	
Response Relevance	Relevance to the patient’s symptoms and preceding context
Overall Quality	Overall relevance, clarity, naturalness, and usefulness
Premature Assessment Avoidance	No diagnosis or treatment planning before sufficient information
Appropriate Uncertainty	Appropriate expression of uncertainty and need for confirmation
Understandability	Clear, patient-friendly questions and explanations
Unsafe Medical Advice Avoidance	Avoidance of potentially harmful medical advice
Unsupported Diagnosis Avoidance	Avoidance of unsupported disease or symptom-cause claims

same amount and format of preference data. Their principal difference is therefore whether the chosen responses are selected for compatibility with the modeled dialogue progression or sampled randomly within a coarsely matched topic category. We use a DPO learning rate of 5×10^{-6} and $\beta = 0.1$. LoRA uses rank $r = 8$, $\alpha = 16$, and dropout of 0.05, and training runs for one epoch. These settings are shared across all three corpus pairs and all three DPO-trained conditions.

Table 3 summarizes the seven oracle-evaluation axes used in each experimental setting.

Oracle evaluation axes for ESConv: For ESConv–DailyDialog, the axes capture the response strategies and progression of emotional-support dialogue. In addition to the strength and tone of ESConv support behavior, they assess maintenance of the supporter role, a non-directive response strategy, alignment between response strategy and dialogue stage, the timing of advice, and naturalness as Japanese dialogue [6, 9, 19].

Oracle evaluation axes for MathDial: For MathDial–WildChat-1M, the axes evaluate not only mathematical correctness, but also whether a response encourages the learner’s own reasoning and provides pedagogical support appropriate to the learner’s current understanding [8]. They assess whether the learner has room to think, whether the system accurately diagnoses the learner’s reasoning and locates errors, and whether it selects among questions, hints, explanations, and answer revelation according to the dialogue stage.

Oracle evaluation axes for MediTOD: For MediTOD–WildChat-1M, the axes capture the appropriateness of medical history-taking, response quality, and medical safety [14]. They assess relevance to the patient’s preceding symptoms

Table 4. Human-evaluation items for each experiment (seven-point scale; 1 = not at all applicable, 7 = very applicable)

Item	Question
Corpus pair: ESConv × DailyDialog	
Q1	Overall, this is a good response for supporting the help seeker.
Q2	The response acknowledges the help seeker’s feelings and speaks gently.
Q3	The response maintains a supportive stance that is consistent with the preceding conversation.
Q4	The response seeks to understand the person without imposing instructions or conclusions.
Q5	The timing of advice or suggestions is appropriate for this stage of the conversation.
Q6	The response fits well with what has been discussed so far.
Q7	The response is natural and easy to read as dialogue.
Final choice	Select the response that best acknowledges distress and supports continued conversation.
Corpus pair: MathDial × WildChat-1M	
Q1	The response leaves room for the learner to think, explain, and try a solution independently.
Q2	The response accurately identifies whether the learner’s reasoning is correct and where it is wrong or incomplete.
Q3	The response specifically narrows down what the learner should review or consider next.
Q4	The questions, hints, or explanations are accurate and match the learner’s current understanding.
Q5	The learner can tell what to think about, calculate, or check next.
Q6	The amount and timing of the answer or solution revealed are appropriate for the current stage of the dialogue.
Q7	The use of questions, hints, explanations, and comprehension checks is appropriate for the dialogue stage.
Final choice	Select the tutoring response that best helps the learner progress independently.
Corpus pair: MediTOD × WildChat-1M	
Q1	The response gathers necessary new information without repeating questions already answered.
Q2	The response does not rush to a diagnosis or action while information is still insufficient.
Q3	The response does not assert what remains unknown and appropriately indicates what still needs to be confirmed.
Q4	The response does not strongly prescribe medication, treatment, or care-seeking when information is insufficient.
Q5	The response does not assert a diagnosis or cause of symptoms without sufficient grounds.
Final choice	Select the clinician response that best gathers information and avoids premature conclusions.

and context, overall response quality, avoidance of premature assessment, appropriate uncertainty, understandability to the patient, and avoidance of unsupported diagnosis or unsafe medical advice. Unsafe Medical Advice Avoidance and Unsupported Diagnosis Avoidance measure how well a response avoids unsafe advice or unsupported diagnosis, rather than how often it produces them. As with the other axes, higher scores therefore indicate more desirable responses.

Human-evaluation items: The human evaluation uses questions that make the response strategies required in each dialogue setting readily interpretable to raters. For every item, 1 means “not at all applicable” and 7 means “very applicable.” Participants were not told the model names or training procedures and judged only the responses.

The final-choice options were common to all domains: Response A, Response B, Response C, Almost the same, and Cannot judge.

4.4 Results

This section presents the human- and oracle-evaluation results for each experimental setting.

Experimental Setting 1: ESConv × DailyDialog

Figure 3 presents the human-evaluation results for Experimental Setting 1. BASiS had the highest mean on all seven axes and significantly exceeded both Base and Random on six, with advice timing as the exception (Holm-adjusted $p < .05$). ESConv Support Behavior Strength averaged 4.22, 5.64, and 3.97 for Base, BASiS, and Random, respectively;




Fig. 3. Human-evaluation results for ESConv-DailyDialog. Bars show means of participant-level averages over 10 contexts; error bars show participant-level bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over three pairs ($*p < .05$, $**p < .01$).




Fig. 4. Oracle-evaluation results for ESConv-DailyDialog. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs ($*p < .05$, $**p < .01$, $***p < .001$).

ESConv Tone Similarity averaged 4.59, 6.17, and 4.43. Kendall's W ranged from .479 to .876 on significant axes. Premature Advice Avoidance also favored BASiS in the mean, but its post-hoc comparisons were nonsignificant. BASiS accounted for 74.4% of final choices.

In the oracle evaluation for Experimental Setting 1 shown in Figure 4, BASiS likewise had the highest mean on all seven axes. For Base, Target-only, BASiS, and Random, ESConv Support Behavior Strength averaged 8.26, 8.22, 8.65, and 8.21, and ESConv Tone Similarity averaged 8.21, 8.25, 8.52, and 8.13; BASiS significantly exceeded all three conditions on both axes. Premature Advice Avoidance averaged 8.57, 8.76, 9.44, and 8.92, again with significant differences from all three conditions. Naturalness was significantly higher than Random, while Target-only did not significantly exceed Base on any axis.

Experimental Setting 2: MathDial $\times$ WildChat-1M

Figure 5 presents the human-evaluation results for Experimental Setting 2. BASiS significantly exceeded both Base and Random on all seven axes (Holm-adjusted $p < .05$). Equitable Tutoring averaged 4.28, 5.42, and 3.73, and Feedback Actionability averaged 4.13, 5.43, and 3.72. Participants therefore favored BASiS both for preserving room for learner reasoning and for clarifying the next action. Kendall's W ranged from .716 to .939. BASiS accounted for 61.1% of final choices.

In the oracle evaluation for Experimental Setting 2 shown in Figure 6, BASiS exceeded all three comparison conditions on all seven axes; all 21 comparisons were significant (all $p < .001$). For Base, Target-only, BASiS, and Random, Equitable Tutoring averaged 6.77, 7.02, 7.75, and 6.00. Learner Reasoning Diagnosis averaged 7.38, 7.76, 8.70, and 6.97, and Mistake Location and Targeting averaged 7.53, 7.74, 8.85, and 7.27. The 1.85-point BASiS–Random difference on Feedback Actionability was the largest among the seven axes. Target-only had a higher mean than Base on all axes, but none of the differences was significant.

Experimental Setting 3: MediTOD $\times$ WildChat-1M




Fig. 5. Human-evaluation results for MathDial-WildChat-1M. Bars and error bars show participant-level means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction ($*p < .05$, $**p < .01$).




Fig. 6. Oracle-evaluation results for MathDial-WildChat-1M. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs (all $***p < .001$).

Figure 7 presents the human-evaluation results for Experimental Setting 3. BASiS had the highest mean on all five axes. It significantly exceeded both Base and Random on Premature Assessment Avoidance, Unsafe Medical Advice Avoidance, and Unsupported Diagnosis Avoidance (Holm-adjusted $p < .05$), with Kendall’s W from .799 to .813. The Friedman tests were nonsignificant for Coverage Without Redundancy and Appropriate Uncertainty. BASiS was selected most often in the final choice, at 44.3%.

In the oracle evaluation for Experimental Setting 3 shown in Figure 8, BASiS likewise had the highest mean on all seven axes. For Base, Target-only, BASiS, and Random, Response Relevance averaged 4.46, 4.70, 5.09, and 4.91, and Overall Quality averaged 5.16, 5.26, 5.69, and 5.46. Appropriate Uncertainty averaged 7.84, 7.96, 8.26, and 7.92, with BASiS significantly exceeding all three conditions. Unsafe Medical Advice Avoidance was significantly higher than Random, and Unsupported Diagnosis Avoidance was significantly higher than Base and Random. Target-only did not significantly exceed Base on any MediTOD axis.

Table 5 summarizes the final-choice distribution. BASiS was selected most often in every domain: 74.4% for ESConv, 61.1% for MathDial, and 44.3% for MediTOD. In MediTOD, this share was not a majority but exceeded Random by 17.2 percentage points.

4.5 Discussion

We evaluated BASiS in three distinct dialogue settings: emotional support, mathematics tutoring, and medical history-taking. BASiS had the highest mean on all 19 human-rated axes, significantly exceeded Base and Random on 16 axes each, and was the most frequent final choice in every domain. It also had the highest mean on all 21 oracle axes and significantly exceeded Base, Target-only, and Random on 18, 13, and 15 axes, respectively. We next discuss the strengths of BASiS revealed by these results and the reasons for the observed improvements.

Discussion 1: Reproducing support behavior in emotional-support dialogue

Human raters preferred BASiS over Base and Random on support behavior strength, an empathetic and accepting tone, supporter-role consistency, non-directiveness, contextual fit, and naturalness; BASiS also received 74.4% of final




Fig. 7. Human-evaluation results for MediTOD-WildChat-1M. Bars and error bars show means for the seven complete participants and participant-level bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction ($*p < .05$).




Fig. 8. Oracle-evaluation results for MediTOD-WildChat-1M. Bars show condition means and bootstrap 95% confidence intervals. Brackets show significant BASiS comparisons after Holm correction over all six pairs ($*p < .05$, $**p < .01$).

Table 5. Counts and percentages of final choices in the human evaluation

Choice	ESConv	MathDial	MediTOD
Base	14 (15.6%)	16 (17.8%)	11 (15.7%)
BASiS	67 (74.4%)	55 (61.1%)	31 (44.3%)
Random	6 (6.7%)	12 (13.3%)	19 (27.1%)
Almost the same	2 (2.2%)	3 (3.3%)	8 (11.4%)
Cannot judge	1 (1.1%)	4 (4.4%)	1 (1.4%)

choices. Because both local expression and the broader stance of acknowledging feelings without imposing a solution were rated highly, BASiS appears to reproduce ESConv’s supporter behavior as an integrated response strategy. One explanation is that selecting data compatible with observation labels such as `reflect_validate` and transitions toward empathetic support teaches the model to maintain a supportive role, rather than merely increasing the use of gentle wording. Premature Advice Avoidance also had the highest human mean under BASiS, while the oracle evaluation showed significant differences from all three comparison conditions. Although the human post-hoc comparisons were nonsignificant, the mean trends agree across the two evaluations. Deferring advice therefore appears as a characteristic of BASiS that could be measured more precisely with additional dialogue contexts.

Discussion 2: Reproducing pedagogical response strategies in mathematics tutoring

Human raters preferred BASiS over Base and Random on all seven axes, including encouraging learner reasoning, diagnosing reasoning, locating mistakes, providing useful support, calibrating answer revelation, and matching teacher moves to the dialogue stage. BASiS also received 61.1% of final choices, indicating that participants regarded the responses as desirable mathematics tutoring overall. The oracle evaluation likewise found significant differences from all three comparison conditions on every axis, making this the most consistent pattern among the three domains. In particular, gains in Feedback Actionability indicate that BASiS did not merely generate correct content, but helped

learners identify what to think about, calculate, or verify next. BASiS selects examples whose transitions from the learner’s inferred state to questions, hints, or explanations match those in MathDial. This selection likely exposed the model to pedagogical scaffolding that moves a learner to the next reasoning step, explaining the broad gains from Equitable Tutoring through Teacher Move–Stage Alignment.

Discussion 3: Response quality and safety in medical history-taking

Human raters preferred BASiS over Base and Random on avoiding premature assessment, unsafe medical advice, and unsupported diagnosis. BASiS was also the most frequent final choice at 44.3%, 17.2 percentage points above Random. Its clearest strength in medical history-taking therefore appears to be a cautious response strategy that withholds conclusions until sufficient information is available. By evaluating candidates through state transitions, BASiS can select responses that gather confirmation before moving to assessment, rather than responses that merely contain medically related words. Coverage Without Redundancy and Appropriate Uncertainty did not show significant human condition differences, although BASiS had the highest mean on both. The oracle evaluation also placed BASiS highest on all seven axes and found a significant advantage over all three conditions on Appropriate Uncertainty. Together, these results suggest that the benefits are clearest for safety-critical forms of restraint, while the efficiency of information gathering and the expression of uncertainty may vary more with the dialogue context. More varied symptoms and dialogue stages would allow these aspects to be examined in greater detail.

Discussion 4: Selecting dialogue progression rather than response content alone

Across all three domains, the strongest results concerned not only topic or tone, but also response timing and dialogue progression. Human raters valued acknowledging feelings in emotional support, scaffolding the learner’s next step in mathematics tutoring, and avoiding premature conclusions in medical history-taking. Each response strategy requires the model to consider both the preceding history and the current dialogue stage. This pattern accords with Equation (3). BASiS does not score candidates solely by generic single-turn quality; it uses posterior probabilities based on state transitions and observation labels. Consequently, knowledge about *when* to offer a given kind of response is reflected in selection and appears in generated responses that humans judge desirable across the dialogue as a whole.

Discussion 5: Contributions of the target corpus and data selection

In the oracle evaluation, Target-only did not significantly exceed Base on any of the 21 axes, whereas BASiS exceeded Target-only on 13. A small target corpus can therefore be used more effectively by extracting its dialogue structure and treating that structure as a selection criterion for a large corpus, rather than using it only as direct training data. BASiS also had a higher oracle mean than Random on all 21 axes and exceeded Random on 16 of the 19 human axes. Both conditions used the same amount of training data and the same training configuration, and Random was already restricted to a coarsely matched topic category. The consistent differences therefore indicate that adaptation to a dialogue setting benefits from selecting examples aligned with conversational states, response strategies, and state transitions in addition to broad topical similarity.

Discussion 6: Applicability across three dialogue settings

We applied the same BASiS pipeline to emotional support, mathematics tutoring, and medical history-taking by changing only the reference and selection corpora. BASiS had the highest mean on every human and oracle axis and was the most frequent final choice in every domain. It can therefore extract different corpus-specific response strategies rather than being limited to a single tone. The fact that the strongest gains differed by domain further illustrates the advantage of inducing states and strategies from each corpus instead of imposing one fixed guideline. Thus, although these findings provide clear initial evidence for the effectiveness of BASiS, generalization to broader user populations requires validation with external raters and more dialogue contexts. Future work should broaden the scope of evaluation

and refine the design of corpus-specific states, strategies, and desirable state sets, thereby extending the strengths of BASiS to a wider range of dialogue applications.

5 Conclusion

We proposed BASiS, a method that transfers dialogue progression from a small, high-quality corpus to an LLM through Bayesian data selection. Under explicit constraints on the number and roles of labels, BASiS uses an LLM to derive conversational states, response strategies, state transitions, and probability estimates from the small corpus, then formalizes them as a Bayesian dialogue model. It uses this model to evaluate candidate responses in a large general-purpose corpus and trains a LoRA adapter on the resulting preference data. This mechanism uses the response strategies and multi-turn progression captured in a small corpus as criteria for selecting training data, without requiring detailed dialogue guidelines to be written manually for each setting. By representing conversational states as a probability distribution and updating that distribution with observations from candidate responses, BASiS selects data according to the dialogue history and the state transition induced by each response.

In comparative experiments across emotional support, mathematics tutoring, and medical history-taking, BASiS achieved the highest mean on all 19 human-rated axes and was the most frequent final choice in every domain. It also achieved the highest mean on all 21 oracle axes. Together, the evaluations suggest that BASiS can incorporate empathetic support, pedagogical guidance suited to learner understanding, and response strategies that avoid unwarranted certainty when information is insufficient. In the oracle evaluation, Target-only did not significantly outperform Base on any axis, and BASiS broadly outperformed Random. These findings demonstrate the value of using the small corpus as a source of interactional knowledge and selecting training data beyond coarse topical similarity.

Future work will evaluate other models, investigate how to design conversational states, response strategies, and desirable state sets for different dialogue domains, and conduct human evaluations with external raters and more dialogue contexts. Through these efforts, we aim to extend BASiS across applications including counseling, education, and medicine.

References

- [1] Lichang Chen, Shiyang Li, Jun Yan, Hai Wang, Kalpa Gunaratna, Vikas Yadav, Zheng Tang, Vijay Srinivasan, Tianyi Zhou, Heng Huang, and Hongxia Jin. 2024. AlpaGasus: Training a Better Alpaca with Fewer Data. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=FdVXgSJhvz>
- [2] Ananya Ganesh, Martha Palmer, and Katharina Kann. 2021. What Would a Teacher Do? Predicting Future Talk Moves. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*. Association for Computational Linguistics, Online, 4739–4751. doi:10.18653/v1/2021.findings-acl.418
- [3] Prakhar Gupta, Yang Liu, Di Jin, Behnam Hedayatnia, Spandana Gella, Sijia Liu, Patrick Lange, Julia Hirschberg, and Dilek Hakkani-Tur. 2023. DialGuide: Aligning Dialogue Model Behavior with Developer Guidelines. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 14031–14047. doi:10.18653/v1/2023.findings-emnlp.935
- [4] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. The Tenth International Conference on Learning Representations (ICLR 2022). <https://openreview.net/forum?id=nZeVKeeFYf9>
- [5] Yanran Li, Hui Su, Xiaoyu Shen, Wenjie Li, Ziqiang Cao, and Shuzi Niu. 2017. DailyDialog: A Manually Labelled Multi-turn Dialogue Dataset. In *Proceedings of the Eighth International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Asian Federation of Natural Language Processing, Taipei, Taiwan, 986–995. <https://aclanthology.org/I17-1099/>
- [6] Siyang Liu, Chujie Zheng, Orianna Demasi, Sahand Sabour, Yu Li, Zhou Yu, Yong Jiang, and Minlie Huang. 2021. Towards Emotional Support Dialog Systems. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Association for Computational Linguistics, Online, 3469–3483. doi:10.18653/v1/2021.acl-long.269
- [7] Wei Liu, Weihao Zeng, Keqing He, Yong Jiang, and Junxian He. 2024. What Makes Good Data for Alignment? A Comprehensive Study of Automatic Data Selection in Instruction Tuning. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=BTKAeLqLMw>

- [8] Jakub Macina, Nico Daheim, Sankalan Pal Chowdhury, Tanmay Sinha, Manu Kapur, Iryna Gurevych, and Mrinmaya Sachan. 2023. MathDial: A Dialogue Tutoring Dataset with Rich Pedagogical Properties Grounded in Math Reasoning Problems. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. Association for Computational Linguistics, Singapore, 5602–5621. doi:10.18653/v1/2023.findings-emnlp.372
- [9] Remi Mir, Bjarke Felbo, Nick Obradovich, and Iyad Rahwan. 2019. Evaluating Style Transfer for Text. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*. Association for Computational Linguistics, Minneapolis, Minnesota, 495–504. doi:10.18653/v1/N19-1049
- [10] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. 2022. Training Language Models to Follow Instructions with Human Feedback. In *Advances in Neural Information Processing Systems*, Vol. 35. Curran Associates, Inc., Red Hook, NY, USA, 27730–27744.
- [11] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 53728–53741. doi:10.52202/075280-2338
- [12] Hannah Rashkin, Eric Michael Smith, Margaret Li, and Y-Lan Boureau. 2019. Towards Empathetic Open-domain Conversation Models: A New Benchmark and Dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*. Association for Computational Linguistics, Florence, Italy, 5370–5381. doi:10.18653/v1/P19-1534
- [13] Eric Rudolph, Philipp Steigerwald, and Jens Albrecht. 2026. Transition-Matrix Regularization for Next Dialogue Act Prediction in Counselling Conversations. In *Findings of the Association for Computational Linguistics: ACL 2026*. Association for Computational Linguistics, San Diego, California, United States, 25442–25457. doi:10.18653/v1/2026.findings-acl.1271
- [14] Vishal Vivek Saley, Goonjan Saha, Rocktim Jyoti Das, Dinesh Raghu, and Mausam. 2024. MediTOD: An English Dialogue Dataset for Medical History Taking with Comprehensive Annotations. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*. Association for Computational Linguistics, Miami, Florida, USA, 16843–16877. doi:10.18653/v1/2024.emnlp-main.936
- [15] Chenwei Wan, Matthieu Labeau, and Chloé Clavel. 2025. EmoDynamiX: Emotional Support Dialogue Strategy Prediction by Modelling MiXed Emotions and Discourse Dynamics. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*. Association for Computational Linguistics, Albuquerque, New Mexico, 1678–1695. doi:10.18653/v1/2025.naacl-long.81
- [16] Xiangyu Wen, Jianyuan Zhong, Zhijian Xu, and Qiang Xu. 2025. Guideline Compliance in Task-Oriented Dialogue: The Chained Prior Approach. In *Findings of the Association for Computational Linguistics: NAACL 2025*. Association for Computational Linguistics, Albuquerque, New Mexico, 6765–6791. doi:10.18653/v1/2025.findings-naacl.377
- [17] Mengzhou Xia, Sadhika Malladi, Suchin Gururangan, Sanjeev Arora, and Danqi Chen. 2024. LESS: Selecting Influential Data for Targeted Instruction Tuning. In *Proceedings of the 41st International Conference on Machine Learning (Proceedings of Machine Learning Research, Vol. 235)*. PMLR, Vienna, Austria, 54104–54132. <https://proceedings.mlr.press/v235/xia24c.html>
- [18] Bolin Zhang, Jiahao Wang, Qianlong Du, Jiajun Zhang, Zhiying Tu, and Dianhui Chu. 2025. A Survey on Data Selection for LLM Instruction Tuning. *Journal of Artificial Intelligence Research* 83, Article 32 (2025), 14 pages. doi:10.1613/jair.1.17625
- [19] Saizheng Zhang, Emily Dinan, Jack Urbanek, Arthur Szlam, Douwe Kiela, and Jason Weston. 2018. Personalizing Dialogue Agents: I have a dog, do you have pets too?. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. Association for Computational Linguistics, Melbourne, Australia, 2204–2213. doi:10.18653/v1/P18-1205
- [20] Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. 2024. WildChat: 1M ChatGPT Interaction Logs in the Wild. The Twelfth International Conference on Learning Representations (ICLR 2024). <https://openreview.net/forum?id=Bl8u7ZRLbM>
- [21] Chunting Zhou, Pengfei Liu, Puxin Xu, Srinu Iyer, Jiao Sun, Yuning Mao, Xuezhe Ma, Avia Efrat, Ping Yu, Lili Yu, Susan Zhang, Gargi Ghosh, Mike Lewis, Luke Zettlemoyer, and Omer Levy. 2023. LIMA: Less Is More for Alignment. In *Advances in Neural Information Processing Systems*, Vol. 36. Curran Associates, Inc., Red Hook, NY, USA, 55006–55021. doi:10.52202/075280-2400